

# Logistic Regression

To get some practice with logistic regression, we'll first be looking at American election data from the National Election Study from 1948 through 2002. Afterward, we'll be taking a look at a classic classification task: email spam filtering.

## A note on caret

In the following, use the `caret` package to obtain cross-validated estimates of  $\alpha$  and  $\lambda$  for regularized logistic regression. You can simply do 5-fold cross-validation without any repeats (instead of 10-fold with 3 repeats). (If you want better predictive power, then do a *second* run of `train()` where you look more finely at values of  $(\alpha, \lambda)$  close to the optimal values you found with a coarser search. However, it's not that important.)

When calling `train()`, you should set `classProbs=TRUE` in the control parameters. Also, set `summaryFunction` to either `twoClassSummary` or `multiClassSummary` depending on whether you want your metric of model quality to be area under the ROC or log loss.

## National Election Study Analysis

If you get stuck, look at section 5.1 of Gelman and Hill as necessary.

For the tasks described below:

- You should use regularized logistic regression in conjunction with the `caret` package.
- Consider adding interaction terms to your fits; if you do so, again use the `caret` package to determine whether or not adding them improves your model.
- Plot the ROC curves for your best models.

- Finally, interpret the coefficients of your models.

Feel free to play around with the data and find interesting relationships.

## Getting started

Refer back to the assignment on factors if you need a refresher on how factors work.

- Use the `read.dta()` function from the `foreign` library to load `elections.dta`.
- Select the columns `year`, `age` through `religion`, `vote`, and `presvote`. Restrict to years with a presidential election.
- If you check the levels of the dataframe's factors (with `lapply(df[sapply(df, is.factor)], levels)`), you'll see that the data needs some cleaning. Modify the `levels()` of each factor so the descriptive text is more concise. You can replace missing values with NAs (using `addNA()`), but note that a missing survey response often carries its own information and corresponds to its own category, and as such should not by default be replaced with a NA and subsequently imputed.
- Replace NAs with column means if the column they're in is numeric. If the column is a factor, instead replace each NA with a randomly chosen level of the factor such that the proportion of each factor level to the number of entries stays unchanged after replacement.
- You'll want to expand each factor out into *indicator variables*. You can use the function you previously wrote in the assignment on factors to do so. Alternatively, you can learn how to use the [dummies package](#) or how to use `model.matrix` to do so.

## Exploring the data

Before you do any data analysis, it's typically a good idea to do some basic visualizations and get a sense of what you're working with.

Use R's `mosaicplot()` function to make a couple mosaic plots from the cleaned and simplified dataset. For example, try:

```
mosaicplot(table(df$income, df$presvote))
```

## Things to predict

- Predict support for George H. W. Bush in the 1992 election. (Restrict consideration to people who actually voted!)

- Predict party support for different years and look at how the coefficients of the predictors change over time.
- For voters who didn't vote, predict how they *would have* voted. If you aggregate these predictions by election year, how do they appear to change over time?

## The CSDMC2010 SPAM Corpus

You'll be looking at the data from the [CSDMC2010 SPAM Corpus](#). The data is in `spam-emails.csv`, with `spam-emails-key.txt` giving the correct classification as spam or not-spam.

- Use the [tm package](#) to construct a `DocumentTermMatrix` from the data, where each row represents a document, each column represents a word, and each entry contains the word frequency for the associated word in the associated document.
- Use the `caret` package to find the optimal values of  $(\lambda, \alpha)$  for regularized logistic regression in the classification of the documents.
- Using the optimal hyperparameters you found, train a regularized logistic regression model on the whole dataset. Look at the ROC curve.

Some things to keep in mind (read these before you begin):

- For a reference about text preprocessing with `tm`, look [here](#) or [here](#).
- Not every single email in the dataset successfully made it into `spam-emails.csv`. You'll want to do an `inner_join()` (from `dplyr`) on the dataset and the classification key, retaining only the rows which are successfully matched.
- Remove columns corresponding to words that show up fewer than 10 times in total throughout the entire corpus.
- The matrix of documents and word frequencies is very sparse – don't scale it! If you do so, it will make it non-sparse, which is bad!