

Nonlinear Methods: Regression and Classification

Signal Data Science

In this lesson, we will explore a collection of standard nonlinear regression and classification techniques. For regression, we'll be predicting wine quality in terms of chemical properties, and for classification, we'll be `TODO` At the end, we'll combine all of these methods with [stacking](#) to create a model which performs better than any individual technique alone.

Getting started

First, we'll load and visualize both datasets.

The wine quality dataset can be found in the wine-quality dataset folder (with documentation in `winequality.names.txt`) or downloaded on the [UCI Datasets page](#). The white wine dataset is three times the size of the red wine dataset, so we'll be focusing on the former – with more data and consequently a greater “resolution”, the advantages of nonlinear techniques become more apparent.

- Load the white wine dataset from `winequality-white.csv`. Read the associated documentation. Replace each space in the column names with an underscore (“_”).
- Use `qplot(...) + geom_smooth()` to plot wine quality against each of the individual variables representing chemical properties in the white wine dataset. Which variables are strongly *and* nonlinearly associated with wine quality?

(todo)

Elastic net regularization

Before using nonlinear methods to predict white wine quality, we will use elastic net regularized linear and logistic regression to calculate a performance

“baseline” to which we can compare the performance of nonlinear methods.

First, we’ll write some utility functions for using `caret`’s `train()`, which will serve as a standardized interface for applying a variety of different nonlinear techniques. `caret` supports [a very large number of different models](#)!

- Write two functions `caret_reg(x, y, method, grid)` and `caret_logit(x, y, method, grid)` which return the output of calling `caret`’s `train()` for a regression task and a binary classification task respectively with predictors `x`, a target variable `y`, a model type `method`, and a grid of hyperparameters `grid`. Follow these specifications:
 - The functions should both set the seed to 1.
 - The `trControl` parameter of `train()` should be set for repeated cross-validation with 1 repeat and 3 folds. Set `verboseIter=TRUE` as well.
 - For classification, the control parameters should include `classProbs=TRUE` and `summaryFunction=twoClassSummary`.
 - In the call to `train()`, set `metric="RMSE"` for regression and `metric="ROC"` (area under the ROC curve) for classification. For both, pass in the parameter `preProcess=c("center", "scale")`.

Multivariate adaptive regression splines

Multivariate adaptive regression splines (MARS) is an extension of linear models that uses *hinge functions*. It models a target variable as being a linear combination of functions of the form $\max(0, \pm(x_i - c))$ where x_i can be any of the predictors in the dataset.

- Look at the pictures on the [Wikipedia page for MARS](#) to get some intuition for how MARS works.

By increasing the *degree* of a MARS model, one can allow for *products* of multiple hinge functions (e.g., $\max(0, x_1 - 10) \times \max(0, 2 - x_3)$), which models interactions between the predictor variables.

Intuitively, one can think of a degree 1 MARS model with p predictor variables as being a piecewise linear combination of hyperplanes – with 1 predictor variable you’re [pasting different lines together](#), with 2 predictor variables you’re [pasting planes together](#), and so on and so forth. Raising the degree then allows more complicated nonlinear interactions to show up.

MARS is implemented as `earth()` in the `earth` package and can be used with `train()` by setting `method="earth"`. It has two hyperparameters to tune, `degree` and `nprune`; the `nprune` parameter is the maximum number of additive terms allowed in the final model (so it controls model complexity).

- Use `caret`'s `train()` to fit a MARS model for white wine quality. Use a grid search to find the optimal hyperparameters, trying `degree=1:5` and `nprune=10:20`.
- Compare the RMSE of the optimal MARS model with the previously obtained RMSEs for white wine quality.
- Pass the optimal hyperparameters into `earth()` directly and examine the resulting model (with `print()` and `summary()`). Interpret the results.
 - When looking at the model, `h(...)` represents a term of the form $\max(0, \dots)$.

Although MARS doesn't usually give results as good as those of more complicated techniques, MARS models are easy to fit and interpret while being more flexible than just a simple linear regression. Degree 1 models can also be built *very* rapidly even for large datasets.

K-Nearest Neighbors

K-Nearest Neighbors (KNN) is one of the simplest possible nonlinear regression techniques.

First, we pick a value of k . Next, suppose that we have a dataset of n points, where each \mathbf{x}_i is associated with a target variable taking on value y_i . Finally, suppose that we have a new point \mathbf{x}^* and we want to predict the associated value of the target variable. To do so, we find the k points \mathbf{x}_i which are closest to \mathbf{x}^* , look at the associated values of y_i , and take their average. That's all!

KNN is implemented in R as `kknn()` in the `kknn` package. It can be used with `caret`'s `train()` by setting `method="kknn"`. There's just a single hyperparameter to tune – the value of k . A larger value of k helps guard against overfitting, but will make the model less sensitive to fine-grained structure in the data.¹

- Use `caret` to train a k -NN model for white wine quality using `tuneLength=10`. Compare the minimum RMSE obtained to the RMSE for a regularized linear model.

In general, we can get better predictions by using information about what happens at a greater distance from the point of interest.

¹One way to interpret k -NN is that it's equivalent to the imposition of a Bayesian prior saying that your dataset is sufficiently fine-grained enough that the value of the target variable at any given point is completely determined by the value of the target variable at nearby points. Given enough granularity, this is in some sense guaranteed to be true (assuming your target variable is a reasonably smooth function of the feature space), but often it's not *perfectly* true. Later on, we'll be looking at support vector machines with a Gaussian basis / radial kernel function, which can be interpreted as a sort of "regularized k -nearest neighbors model" and performs very well in practice for classification.

Regression tree models

We'll proceed to explore three different types of nonlinear techniques: regression trees, random forests, and gradient boosted trees. All three of these can be used for both regression and classification.

Regression trees are the simplest and very easily interpretable, but their performance is often poor and they tend to overfit. We can train an *ensemble* of regression trees and combine them together into a *random forest*, or we can keep training regression trees in an iterative manner to keep improving a single model in a technique known as *boosting*.

Using regression trees

You'll need the `rpart()` function in the `rpart` package to construct regression tree models. It's used in the same way as `lm()`, both in how the model is constructed and in how predictions are made using the model.

- For both red and white wines, create regression tree models predicting wine quality with the other features in the dataset. View them (by just calling `print()` on the models) and interpret the differences between the two models.
- Compare the regression tree model for white wine quality with the previously obtained MARS model for white wine quality.

There is a single hyperparameter involved in the fitting of a regression tree: the *complexity parameter*, usually denoted `cp`. (It defaults to `0.01` if not explicitly specified in the call to `rpart()`.) As we grow a regression tree, we only make another 'split' in the tree if the associated incremental increase in the overall R-squared is greater than the value of `cp`. A higher value of `cp` helps us guard against overfitting to the data, but it can also stop us from growing the tree to a sufficient depth.

As before, we can use `caret`'s `train()` to test different values of `cp`. Since there's only a single hyperparameter to optimize, we can again use the `tuneLength=10` parameter.

- Use the `caret` package to fit a regression tree for the prediction of white wine quality with its chemical characteristics. Compare the RMSE value for the best fit with the RMSE from regularized linear regression and *k*-NN.

Using random forests

Next, you'll be using `randomForest()` from the [randomForest](#) package, which is a more sophisticated nonlinear regression and classification technique.

Theoretical overview

In short, a *random forest* trains a lot of different regression trees and averages them together, with these two conditions on the regression trees:

1. Each regression is trained on a subset of the original data which is sampled *with replacement*. This technique is known as *bagging* and helps combat overfitting. (Usually as many samples are drawn as there are data points in the original dataset.)
2. At each split of each regression tree, only a random subset of the original predictors are considered as candidate variables for the split (usually \sqrt{p} candidate predictors for p total predictors). This prevents very strong predictors from dominating certain splits and thereby *decorrelates* the regression trees from each other. The size of this random subset, denoted as `mtry`, is the sole hyperparameter needed to fit a random forest model.

As a bonus, we can fit each data point in the training data to the trees that *weren't* trained on that data point (and average the subsequent predictions) to obtain an *out-of-bag error*, which is an estimate for the generalizable error of our model. With this, we don't really have to use cross-validation to estimate the generalizable error of our model.

- Read [Edwin Chen's Quora answer](#) on how random forests work.

Random forests in R

In general, the `randomForest()` function is used in a manner analogous to `rpart()` and `lm()`. There are two details to pay attention to:

First, it's important to pay some attention to the choice of the `mtry` hyperparameter. It's usually advised to try either `mtry = floor(sqrt(p))` or `mtry = floor(p/3)` (for a dataset with p predictors); the former should be used when $p/3$ rounds to 1-2 or when we don't have very many predictors relative to the number of data points ($p \ll n$), and the latter should be used otherwise. Also, it's *always* wise to try `mtry = p` if possible.² The computation time re-

²Setting the value of `mtry` carefully is of [debatable importance](#). [Random Forests for Classification in Ecology](#) by Cutler *et al.* (2007) reports that performance isn't very sensitive to `mtry`, whereas [Conditional variable importance for random forests](#) by Strobl *et al.* (2008) reports the opposite. Finally, [Random Forests: some methodological insights](#) by Genuer *et al.* (2008) finds varying importance for `mtry` depending on properties of the dataset. **All considered**, I think it's fine to initially just try $p/3$, \sqrt{p} , and p , and to decide if further tuning is warranted based on those results.

quired for larger datasets can be quite significant, so `mtry = floor(sqrt(p))` is fine if all you want is a baseline for nonlinear regression performance; even `mtry = floor(p/3)` can be much slower, both for model fitting and for making predictions.

Second, when using the `predict()` function on a random forest model, there is an [important point](#) to keep in mind. Suppose that we've run `rf = randomForest(y ~ x, df)` and we want to evaluate the RMSE associated with that fit. To that end, we'd like to generate predictions on the original dataset. We can run one of two commands:

1. `predict(rf)`, which will make predictions for each data point only with trees which weren't trained on that data point, thereby allowing us to calculate a generalizable *out-of-bag error*, and
2. `predict(rf, df)`, which will use the *entire tree* and seem to indicate severe problems with overfitting if we calculate the associated RMSE.

Usually, what you want is `predict(rf)`, not `predict(rf, df)`.

Anyway, let's get some practice with random forests:

- With the white wine dataset, fit two random forest models for wine quality as a function of the other variables, setting `mtry = floor(sqrt(p))` and `mtry = p`.
- Make out-of-bag predictions with both of the random forest models and calculate the associated RMSE values. Compare the RMSEs to previously obtained RMSEs.
- Use `caret`'s `train()` function with `method="parRF"` to tune over `mtry=c(3, 5, 7, 11)`. Compare the minimum cross-validated RMSE with the out-of-bag RMSE estimate. (The `parRF` method is a parallelized version of `randomForest()`.)

Using gradient boosted trees

Boosting is a technique that iteratively improves a decision tree ensemble with more and more decision trees.³ Specifically, *gradient boosting* is a very powerful nonlinear technique and is one of the best "off-the-shelf" machine learning models.⁴ They train relatively quickly, they can pick up on fairly complicated nonlinear interactions, you can guard against overfitting by increasing the shrinkage parameter, and their performance is difficult to beat.

³The first implementation of boosting – or at least the most famous – is [AdaBoost](#), which can be considered to be like a [special case of gradient boosting](#), a more modern form of boosting.

⁴See Ben Kuhn's [comments](#) on gradient boosting.

However, they're a little more complicated than random forests; there are more hyperparameters to tune, and it's much more difficult to parallelize gradient boosted trees.⁵

Intuitively, one can think of boosting as iteratively improving a regression tree ensemble by repeatedly training a new regression tree on the *residuals* of the ensemble (when making predictions on the dataset) and then incorporating that regression tree into the ensemble.

Gradient boosted trees are implemented in R's `gbm` package as the `gbm()` function. They're also compatible with `caret`'s `train()` – just set `method="gbm"`.

- Use `train()` to perform a *grid search* to optimize the hyperparameters for a gradient boosted tree model (predicting white wine quality from chemical properties).
 - Instead of passing in the `tuneLength` parameter like earlier, use `expand.grid()` to create a grid with `n.trees` set to 500, `shrinkage` set to $10^{\text{seq}(-3, 0, 1)}$, `interaction.depth` set to 1:3, and `n.minobsinnode` set to `seq(10, 50, 10)`.
- With the optimal values of the hyperparameters determined with `train()`, run `train()` again and tune only the value of `n.tree`, trying values from 500 to 5000 in steps of 500. Compare the minimum RMSE to previously obtained RMSEs for other models.

Cubist

Cubist is a nonlinear *regression* algorithm developed by Ross Quinlan with a [proprietary implementation](#). (The single-threaded code is open source and has been [ported to R](#).)

In practice, *Cubist* performs approximately as well as a gradient boosted tree (as far as predictive power is concerned).⁶ Having only two hyperparameters to tune, *Cubist* is a little simpler to use, and the hyperparameters themselves are very easily interpretable.

Broadly speaking, *Cubist* works by creating a *tree of linear models*, where the final linear models are *smoothed* by the intermediate models earlier in the tree. It's usually referred to as a *rule-based model*.

- *Cubist* incorporates a *boosting-like scheme* of iterative model improvement where the residuals of the ensemble model are taken into account when

⁵See [StackExchange](#) for a brief overview of tuning `gbm()` hyperparameters.

⁶In [Subpixel Urban Land Cover Estimation](#) by Walton (2008), *Cubist*, random forests, and support vector regression are compared for a prediction task, and *Cubist* is found to be superior to gradient boosted trees. A [comment on a Ben Kuhn post](#) reports the same result.

training a new tree. Cubist calls its trees *committees*, and the number of committees is a hyperparameter which must be tuned.

- Cubist can also adjust its final predictions using a more complex version of k -NN. When Cubist is finished building a rule-based model, Cubist can make predictions on the training set; subsequently, when trying to make a prediction for a new point, it can incorporate the predictions of the K nearest points in the training set into the new prediction.

As such, there are two hyperparameters to tune, called *committees* and *neighbors*. *committees* is the number of boosting iterations, and the functionality of *neighbors* is easily intuitively understandable as a more complex version of k -NN. The Cubist algorithm is available as `cubist()` in the Cubist package and can be used with `train()` by setting `method="cubist"`.

- Use `caret`'s `train()` to fit a Cubist model for white wine quality. Use a grid search to find the optimal hyperparameter combination, searching over `committees=seq(10, 30, 5)` and `neighbors=0:9`.
- Compare the RMSE of the best Cubist fit with previously obtained RMSEs, particularly the RMSE corresponding to a gradient boosted tree.

Note that Cubist can only be used for *regression*, not for *classification*. Quinlan also developed the [C5.0 algorithm](#), which is for classification instead of regression.

Stacking

[Stacking](#) is a technique in which multiple different learning algorithms are trained and then *combined* together into an ensemble.⁷ The final 'stack' is very computationally expensive to compute, but usually performs better than any of the individual models used to create it.

Ensemble stacking using a `caret`-based interface is implemented in the [caretEnsemble package](#). We'll start off by illustrating how to combine (1) MARS, (2) K-Nearest Neighbors, and (3) regression trees into a stack.

We'll first have to specify which methods we're using and the control parameters:

```
ensemble_methods = c('glmnet', 'knn', 'rpart')
ensemble_control = trainControl(method="repeatedcv", repeats=1,
                                number=3, verboseIter=TRUE,
                                savePredictions="final")
```

Next, we have to specify the tuning parameters for all three methods:

⁷The canonical paper on stacking is [Stacked Generalization](#) by Wolpert (1992).


```
ensemble_tunes = list(
  glmnet=caretModelSpec(method='glmnet', tuneLength=10),
  kknn=caretModelSpec(method='kknn', tuneLength=10),
  rpart=caretModelSpec(method='rpart', tuneLength=10)
)
```

We then create a list of `train()` fits using the `caretList()` function:

```
ensemble_fits = caretList(quality ~ ., df_whitewine,
  trControl=ensemble_control,
  methodList=ensemble_methods,
  tuneList=ensemble_tunes)
```

Finally, we can find the best *linear combination* of our many models by calling `caretEnsemble()` on our list of models:

```
fit_ensemble = caretEnsemble(ensemble_fits)
print(fit_ensemble)
summary(fit_ensemble)
```

By combining three simple methods, we've managed to get a cross-validated RMSE lower than the RMSE for any of the three individual models!

- How much lower does the RMSE get if you add in gradient boosted trees to the ensemble model? (The `caretModelSpec()` function can take a `tuneGrid` parameter instead of `tuneLength`.)

In the [caretEnsemble documentation](#), read about how to use `caretStack()` to make a more sophisticated *nonlinear ensemble* from `ensemble_fits`.

- If you use a gradient boosted tree for `caretStack()`, is it any better than the simple linear combination?
- If you use all of the techniques you've just learned about in a big ensemble, how low can the RMSE get? (This might take a lot of computation time, so it's **optional**, but is fun to look at.)

Closing notes

By now, you've tried a fairly wide variety of nonlinear fitting techniques and gotten some sense for how each of them works. *In practice*, people usually use tree-based methods, especially random forests and gradient boosted trees – they tend to be fairly easily tuned and robust to overfitting. However, it's useful to have a broader overview of the field as a whole.

Also, there are a lot of peculiarities to the interfaces of different nonlinear techniques – when comparing them, `caret` offers a very well-designed interface for all of them, so it's nice to stick to using `train()` and other `caret` methods when possible.

Hyperparameter optimization

You may have noticed that tuning hyperparameters is a very big part of fitting nonlinear methods well! As the techniques become more complex, the number of hyperparameters to tune can grow significantly. Grid search is fine for ordinary usage, but in very complicated situations (10-20+ hyperparameters) it's better to use [random search](#) – otherwise there would just be far too many hyperparameter combinations to evaluate!⁸

- Read the first 4 paragraphs of the [caret package's documentation on random hyperparameter search](#).

The `caret` package is very well-designed, and grid search will usually suffice for your purposes, especially because of its internal optimizations. It's good to be aware that alternatives to grid search exist.

Which model to use?

When trying to do predictive regression modeling, it's usually advised to start out with random forests or gradient boosted trees because they're fairly well understood and perform well out of the box with fairly straightforward tuning.⁹ Random forests are simpler than gradient boosted trees, but both are much simpler than, say, a deep neural net.

For fast parallelized gradient boosted trees in R, use the [xgboost package](#) – it's currently the state of the art. For random trees, the currently best implementation can be used by setting `method="parRF"` in `caret's train()`, which is a parallelized combination of the `randomForest`, `e1071`, and `foreach` packages.

In the future, you'll learn about more complex nonlinear regression techniques, which can either be used on their own or be combined with the techniques you've already learned in a larger ensemble. However, defaulting to either random forests or gradient boosted trees works quite well in practice if you want to get a sense of how much predictive improvement you can get from using a nonlinear method.

⁸If you have, say, 15 hyperparameters, even the simplest possible grid search that selects one of two possible values for each hyperparameter still has 2^{15} configurations to iterate over. That will almost assuredly take far too long.

⁹One of the only good comparisons of nonlinear regression techniques is in [BART: Bayesian Additive Regression Trees](#) by Chipman *et al.* (2010), which gives the following ordering (from better to worse): BART, 1-layer neural nets, gradient boosted trees, random forests. Cubist isn't used very much, mostly because almost nobody really knows what it does, even if its results are pretty good in practice. See also [Performance Analysis of Some Machine Learning Algorithms for Regression Under Varying Spatial Autocorrelation](#) by Santibanez *et al.* (2015).