Simple Linear Regression

Next, we'll continue exploring the effects of public education spending.

Massachusetts Test Score dataset

- Load the `MCAS` dataset from the `Ecdat` package into a variable `df` and read about it using `help(MCAS)`.

Cleaning the dataset

- Find the total number of rows.
- Remove the rows with missing values, and compute the number of rows of the resulting data frame.
  - Is the number of rows appreciably smaller?

Simply removing all the rows with even a single missing value instead of filling in the missing values somehow can lead to statistical problems with our analyses. The process of filling in the missing values is called imputation, which we'll cover in greater depth in the future.

Preliminary analysis

We'll start out with some more simple linear regressions before moving to a slightly more advanced technique.

- Compute the correlations of `totsc4` and `totsc8` with the other numeric features in the dataset.[1]
  - Remember to select only the numeric columns when passing in the dataset to `cor()`.
- Form a new data frame `df1` by removing the non-numeric columns (including factors) and `totsc4`.
- Run a regression of `totsc8` against total expenditure per student, `totday`.
  - What does this say about the effect of spending per student on standardized test scores?
- Run a regression of totsc8 against the other features using `lm(totsc8 ~ . , df1)`.

---

[1]You might want to consider why the correlations with `totsc8` are larger than the ones with `totsc4`.

– Should we be including `code` as a predictor? If not, remove it and see how the results change.

- Run a regression of `totsc8` against the 3 predictors with p-value < 0.01 in the above regression.

    – Is the predictive power appreciably lower?

Stepwise linear regression

In general, the problem of *feature selection* is a difficult one. We ideally want to maximize predictive power using as few features as possible, because adding redundant features actually works *against* interpretability and pulls weight away from the less-redundant ones; however, with n features, we have 2^n possible combinations of features to regress against.

The most simplistic way of solving this problem is with stepwise linear regression. In *forward*[^fwd stepwise linear regression, we initialize the linear model with no predictors (so the model is just a constant), and we keep adding predictors which, when added, provide the greatest incremental boost to the model quality. When we reach a point where the incremental improvements are minimal, we stop.[2]

Before learning how to run a stepwise regression in R, briefly read about its implementation.

Use stepwise regression by writing:

```
model_init = lm(totsc8 ~ 1, df1);
model = formula(lm(totsc8 ~ ., df1))
step_reg = step(model_init, model, direction = "forward")
```

- How does the resulting model differ from the model above?

- Interpret the order in which coefficients are added and removed from the stepwise model.

- What does this say about the effect of educational expenditure on student test scores for schools in the dataset?

    – Reconcile this with your earlier results.

- Repeat the above with `totsc8` replaced by `totsc4` and compare the results.

---

[2]The evaluation of "model quality" is often done via the Akaike information criterion, which can intuitively be thought of as being analogous to the *entropy* of a model. (Minimizing the AIC is broadly equivalent to maximizing the entropy in a thermodynamic system.)