

Notes: Alternating Least Squares

Signal Data Science

This document is an exposition of the method of alternating least squares (ALS) for imputation of missing values. The reference is Hastie *et al.* (2014), [Matrix Completion and Low-Rank SVD via Fast Alternating Least Squares](#); ALS is implemented in the `softImpute` package in R.

Theory

The operation of [matrix multiplication](#) allows us to *multiply* two matrices and form a new matrix. It is illustrated below:

$$\begin{array}{c} 4 \times 2 \text{ matrix} \\ \begin{bmatrix} a_{11} & a_{12} \\ \cdot & \cdot \\ a_{31} & a_{32} \\ \cdot & \cdot \end{bmatrix} \end{array} \begin{array}{c} 2 \times 3 \text{ matrix} \\ \begin{bmatrix} \cdot & b_{12} & b_{13} \\ \cdot & b_{22} & b_{23} \end{bmatrix} \end{array} = \begin{array}{c} 4 \times 3 \text{ matrix} \\ \begin{bmatrix} \cdot & x_{12} & x_{13} \\ \cdot & \cdot & \cdot \\ \cdot & x_{32} & x_{33} \\ \cdot & \cdot & \cdot \end{bmatrix} \end{array}$$

Figure 1: An illustration of matrix multiplication, where $x_{12} = a_{11}b_{12} + a_{12}b_{22}$ and $x_{33} = a_{31}b_{13} + a_{32}b_{23}$.

Note specifically the dimensions of the resulting matrix. If \mathbf{A} is an $n \times p$ matrix and \mathbf{B} is a $p \times m$ matrix, the product \mathbf{AB} will have dimensions $n \times m$.¹ What if p is very small compared to n and m ? We will be able to obtain quite a large matrix just from multiplying together two very narrow matrices (one tall and one wide).

In general, we find that it is possible to [decompose](#) large matrices into the *product* of multiple *smaller* matrices. This is the key behind the method of alternating least squares.

The task at hand is that given a matrix \mathbf{X} with many missing entries, we want to construct a filled-in matrix \mathbf{Z} which *minimizes some loss function*. It turns out that we can write a *regularized* cost function which makes this task straightforward, and also that we can write the solution which minimizes the cost function as

¹If the matrices all have real values, we can write $\mathbf{A} \in \mathbb{R}^{n \times p}$, etc.

$$\mathbf{Z} \approx \mathbf{A}\mathbf{B}^\top$$

for an appropriate choice of a tall matrix \mathbf{A} and a wide matrix \mathbf{B}^\top , where the operator \top denotes the *transpose* of a matrix, (flipping a $n \times m$ matrix so that its dimensions become $m \times n$). Note that for the existing data in \mathbf{X} we simply use that rating data directly in the filled-in matrix \mathbf{Z} instead of the approximated values in $\mathbf{A}\mathbf{B}^\top$ (hence the \approx symbol).

Our imputation method is an indirect one in the sense that instead of *directly* trying to calculate missing values from existing ones, we ask what the optimal filled-in matrix \mathbf{Z} would look like and infer the missing values based on an analysis of \mathbf{Z} . Precisely, we are trying to minimize the differences between the filled-in entries of \mathbf{X} and the corresponding entries of $\mathbf{A}\mathbf{B}^\top$ along with a regularization term controlled by a parameter λ .² Our cost function only considers the matrix entries which correspond to existing data (the filled-in values of \mathbf{X}), but the fashion in which we estimate \mathbf{A} and \mathbf{B} operate on the *entirety* of each matrix. Consequently, the entries of $\mathbf{A}\mathbf{B}^\top$ corresponding to *missing* data in \mathbf{X} serve as rating estimates.

Our task is now simply to estimate the matrices \mathbf{A} and \mathbf{B} . It turns out that the optimal estimates are related via the equation

$$\mathbf{B} = (\mathbf{A}^\top \mathbf{A} + \lambda \mathbf{I})^{-1} \mathbf{A}^\top \mathbf{Z}$$

and vice versa with \mathbf{A} and \mathbf{B} switched, where λ is the regularization parameter and \mathbf{I} is the identity matrix.³ For mathematical reasons, this is actually equivalent to running a regularized⁴ least squares regression for each column of \mathbf{Z} with the columns of \mathbf{A} as predictors, with the coefficient estimates corresponding to entries of \mathbf{B} .⁵

As such, this suggests a strategy for estimating \mathbf{A} and \mathbf{B} . First, we start by initializing \mathbf{A} . Next, we use the regression strategy described above both to

²One can ask the question of why we don't simply estimate $\mathbf{A} = \mathbf{B} = \mathbf{0}$ in the degenerate case of \mathbf{X} having *no* missing values. If that were the result, it would contradict the fact that $\mathbf{A}\mathbf{B}^\top$ should be equal to the soft-thresholded SVD of \mathbf{Z} (presented later in this exposition)! The reason is that although $\mathbf{A}\mathbf{B}^\top$ contains information about the *differences* between \mathbf{X} and \mathbf{Z} , we don't try to impute those differences directly (in which case we might stop immediately if there were no differences whatsoever) but rather *infer* them by trying to bring \mathbf{X} and $\mathbf{A}\mathbf{B}^\top$ closer together.

³The identity matrix is a matrix with 1 on the diagonal and 0 elsewhere. Multiplying it by a different matrix leaves that matrix unchanged.

⁴Specifically, this is equivalent to using [Tikhonov regularized linear regression](#) with Tikhonov matrix $\Gamma = (\lambda \mathbf{I})^{1/2}$. This is also called *ridge regression* and reduces to L^2 regularization in the case where Γ is the identity matrix. We are essentially running a linear regression of each column of \mathbf{Z} with the columns of \mathbf{A} as predictors and getting \mathbf{B} back as the coefficient estimates.

⁵We can check that the dimensions match up. Suppose that $\mathbf{Z} \in \mathbb{R}^{n \times p}$, $\mathbf{A} \in \mathbb{R}^{n \times f}$, and $\mathbf{B} \in \mathbb{R}^{p \times f}$. Then the columns of \mathbf{Z} and \mathbf{A} all have n entries each, and so we can run p different linear regressions (one for each column of \mathbf{Z}) and get out f coefficient estimates each time. We therefore estimate $p \times f$ different coefficient estimates in total, which matches up with the dimensions of \mathbf{B} .

generate predictions for \mathbf{Z} and to generate an estimate for \mathbf{B} . Next, we can switch the places of \mathbf{A} and \mathbf{B} in the above equation and use the same process to update \mathbf{Z} and \mathbf{A} . We repeat in this *alternating* fashion until we achieve convergence.

Soft-thresholded SVD

After running the algorithm described above, `softImpute()` returns the imputed matrix \mathbf{Z} in a *special form*. It turns out that the product \mathbf{AB}^\top is related to \mathbf{Z} in yet another fashion!

Taking a step back: in general, *all* matrices can be decomposed into a product of the form \mathbf{UDV}^\top called the [singular value decomposition](#) (SVD) where \mathbf{D} is a diagonal matrix (the only nonzero entries are on the diagonal). We can compute a *modified* version of the SVD for \mathbf{Z} called the *soft-thresholded SVD* formed by taking \mathbf{D} and shrinking the entries on its diagonal toward 0 by a value λ , setting an entry d_i equal to 0 if $|d_i| \leq \lambda$.⁶ With the modified matrix \mathbf{D}^* , we can compute the soft-thresholded SVD as $S_\lambda(\mathbf{Z}) = \mathbf{UD}^*\mathbf{V}^\top$.

The connection between \mathbf{AB}^\top and \mathbf{Z} lies in the somewhat remarkable relation

$$\mathbf{AB}^\top = S_\lambda(\mathbf{Z})$$

for the optimal estimates of \mathbf{A} and \mathbf{B} .

Indeed, `softImpute()` will return three matrices as `$u`, `$d`, and `$v`, corresponding to the matrices in $S_\lambda(\mathbf{Z}) = \mathbf{UD}^*\mathbf{V}^\top$. From those, we also know \mathbf{AB}^\top , and so the imputed matrix $\mathbf{Z} \approx \mathbf{AB}^\top$ can be calculated.

Dimensionality reduction

From the definition of the soft-thresholded SVD, we see that increasing λ sufficiently high will make every value in \mathbf{D}^* equal to 0. The immediate takeaway is that by calculating the maximum value in \mathbf{D} , we can establish an *upper bound* for the values of λ to test. However, there is a more important and subtler interpretation of the results of ALS in connection with the regularization parameter.

It is likely that the optimal value of λ is one which drives some *but not all* of the values in \mathbf{D} to 0. An $n \times n$ diagonal matrix with a *rank* of k , *i.e.*, k nonzero values on the diagonal can simply be rewritten as a $k \times k$ diagonal matrix without any nonzero values on the diagonal. Our decomposition then becomes the product

⁶Soft-thresholding is basically solving a L^1 regularized cost function very rapidly by looking at the first derivative. Refer back to the theoretical discussion in *Linear Regression: Regularization* for some related details. We can therefore think of soft-thresholded SVD as a sort of L^1 regularized version of SVD which shrinks the singular values closer to 0.

of (1) \mathbf{U} (a tall $n \times f$ matrix), (2) \mathbf{D}^* (a small square $f \times f$ matrix), and (3) \mathbf{V}^\top (a wide $f \times m$ matrix) for some small value of f . We can interpret this as being able to *summarize* both users and movies in terms of f factors, with the columns of \mathbf{U} being factor scores for users and the rows of \mathbf{V}^\top being factor scores for movies.

If a user has factor scores $\mathbf{u} = (u_1, u_2, \dots, u_f)$, a movie has factor scores $\mathbf{m} = (m_1, m_2, \dots, m_f)$, and the diagonal entries of \mathbf{D}^* are given by $\{d_1, d_2, \dots, d_f\}$, then the predicted rating for that user–movie pair is simply given a **weighted inner product** of \mathbf{u} and \mathbf{m} equal to

$$\langle \mathbf{u}, \mathbf{m} \rangle = \mathbf{u}^\top \mathbf{D} \mathbf{m} = \sum_{i=1}^f u_i d_i m_i.$$