



Andrew Ho <kironide@gmail.com>

Week 2 Day 4 (Morning)

3 messages

Jonah Sinick <jsinick@gmail.com>

Thu, Feb 25, 2016 at 2:42 AM

To: Ali Bagherpour <ali.bagherp@gmail.com>, Andrew Ho <Kironide@gmail.com>, Chad Groft <clgroft@gmail.com>, David Bolin <david@bolin.at>, Jacob Pekarek <jpekarek@trinity.edu>, Jaiwithani <jaiwithani@gmail.com>, James Cook <cookjw@gmail.com>, Linchuan Zhang <email.linch@gmail.com>, Matthew Gentzel <magw6270@terpmail.umd.edu>, Olivia Schaefer <taygetea@gmail.com>, Sam Eisenstat <sam.eisenst@gmail.com>, Tom Guo <tomguo4@gmail.com>, Trevor Murphy <trevor.m.murphy@gmail.com>

Today we're going to look at comparison of means between groups.

- Read the Introduction and Test 1 through Test 10 of [100 Statistical Tests by Gopal Kanji](#), writing a function in R to perform each test, with descriptive variable names, in order to make sure that you understand the statements of the tests. Don't worry about memorizing them.
- Read [Calculating p-values](#) (in R).
- Skim [Always use Welch's t-test instead of Student's t-test](#) from "The 20% Statistician."
- Check out the documentation for `pairwise.t.test` in R <https://stat.ethz.ch/R-manual/R-patched/library/stats/html/pairwise.t.test.html>. Read about the Bonferroni Correction https://en.wikipedia.org/wiki/Bonferroni_correction (`pairwise.t.test` uses a modified version).
- Load the dataset "singer" from the "lattice" package. Compute the means by voice part. Apply `pairwise.t.test` to the dataset to determine which differences in height by voice part are statistically significant.
- If you haven't finished *Advanced R*, Chapters 1–6, 10 and 11, continue working through them (including the exercises).
- Take a look at `LWSurveyExploration.zip` (as attached). Explore groupings of the numeric data by category, using `t.test` and `pairwise.t.test` as appropriate. Write up your most interesting finding in a `.Rmd` file.

**LWSurveyExploration.zip**

488K

Jonah Sinick <jsinick@gmail.com>

Thu, Feb 25, 2016 at 10:13 AM

To: Ali Bagherpour <ali.bagherp@gmail.com>, Andrew Ho <Kironide@gmail.com>, Chad Groft <clgroft@gmail.com>, David Bolin <david@bolin.at>, Jacob Pekarek <jpekarek@trinity.edu>, Jaiwithani <jaiwithani@gmail.com>, James Cook <cookjw@gmail.com>, Linchuan Zhang <email.linch@gmail.com>, Matthew Gentzel <magw6270@terpmail.umd.edu>, Olivia Schaefer <taygetea@gmail.com>, Sam Eisenstat <sam.eisenst@gmail.com>, Tom Guo <tomguo4@gmail.com>, Trevor Murphy <trevor.m.murphy@gmail.com>

The case study at the end of the tutorial that <http://www.cyclismo.org/tutorial/R/pValues.html> is from is also good for a real world example using t-tests and calculating p-values: <http://www.cyclismo.org/tutorial/R/cholesterol.html>

[Quoted text hidden]

Jonah Sinick <jsinick@gmail.com>

Thu, Feb 25, 2016 at 3:20 PM

To: Ali Bagherpour <ali.bagherp@gmail.com>, Andrew Ho <Kironide@gmail.com>, Chad Groft

<clgroft@gmail.com>, David Bolin <david@bolin.at>, Jacob Pekarek <jpekarek@trinity.edu>, Jaiwithani <jaiwithani@gmail.com>, James Cook <cookjw@gmail.com>, Linchuan Zhang <email.linch@gmail.com>, Matthew Gentzel <magw6270@terpmail.umd.edu>, Olivia Schaefer <taygetea@gmail.com>, Sam Eisenstat <sam.eisenst@gmail.com>, Tom Guo <tomguo4@gmail.com>, Trevor Murphy <trevor.m.murphy@gmail.com>

I've attached the presentation.

A correction thanks to Ali: The standard deviation of the sampling means is $\text{standardDeviationOfSample} / \sqrt{\text{sampleSize} - 1}$: I was missing the -1.

You can look at the movielens dataset:

<http://grouplens.org/datasets/movielens/> (I used the one with 1,000,000 ratings).

Main question: Which movies are robustly higher rated by women?

Write up your analysis as a .Rmd folder.

Feel free to explore genre information.

Watch out for the [multiple comparisons problem](#)

If you'd like to do something else with the dataset, you can check out [recommenderlab: A Framework for Developing and Testing Recommendation Algorithms](#) ;-).

[Quoted text hidden]

2 attachments



movieGenderDifferences.html
1057K



movieGenderDifferences.Rmd
3K