Web Scraping and NLP

In this lesson, we will focus on natural language processing. Beginning with some classical NLP tasks, we will slowly work our way up to the most complex methods used on text that we scrape directly from the Internet.

Many of these methods will use the *bag-of-words* assumption, where we assume that within each document the *order* of words isn't important. This assumption will only be violated when we use *n-grams*, groups of *n* consecutive words, which *do* preserve information about the order of words.

As you progress through this assignment, write up your findings. At the end, email us with your writeup of your results.

Writing a simple spellcheck function

Spelling correction is one of the most natural and oldest natural language processing tasks. It may seem like a difficult task to you at the moment, but it's surprisingly easy to write a spellchecker that does fairly well. (Of course, companies like Google spend millions of dollars making their spellcheckers better and better, but we'll start with something simpler for now.)

- Read Peter Norvig's How to Write a Spelling Corrector. Recreate it in R and reproduce his results. **After** doing so yourself, read about this 2-line R implementation of Norvig's spellchecker.

Email spam classification

Using a naive Bayes classifier for the task of filtering spam is one of the classic applications of machine learning, going all the way back to 1998. Essentially, the idea is to classify spam probabilistically based on which words appear or don't appear in the email (without taking into account word frequency).

- Read Paul Graham's essay A Plan for Spam, where he describes his improvements to the simplest naive Bayes spam filtration technique which

substantially decreased the false positive rate, as well as Better Bayesian Filtering, where he elaborates further on why his filtration method works better and on some additional improvements he made to it.

Naive Bayes classification

For any given word and email, suppose that W, S, and H are respectively events corresponding to the word appearing in the email, the email being spam, and the email not being spam ("H" for "ham"). From Bayes' theorem, we can say

$$P(S \mid W) = \frac{P(W \mid S)P(S)}{P(W \mid S)P(S) + P(W \mid H)P(H)}.$$

Recent statistics that the probability of any given email being spam is around 80%. However, the simplest Bayesian spam filters assume that there is no *a priori* reason to assume that an incoming email is more likely to be spam and not, and consequently set $P(S) = P(H) = \frac{1}{2}$. We may then simplify the above expression to

$$P(S \mid W) = \frac{P(W \mid S)}{P(W \mid S) + P(W \mid H)}.$$

Intuitively, this is the same as asking: "Of all the occurrences of the word under consideration, what proportion of them appear in spam messages?" This is something we can calculate directly from the training data.

However, how do we combine each evaluation of $P(S \mid W)$ (for each word which appears in the dataset) to find $P(S)$? This is difficult in general, because the appearances of words may be *correlated* – if, say, "birthday" appears in an email, it's more likely than usual that "party" will also appear in the email. Ordinarily, we would have to account for all these "interactions" between different words.

To simplify our task and turn our *full Bayes* classifier into a *naive Bayes* classifier, we assume that whether or not a word appears in an email is *independent* from whether or not any other word appears in the email. This is almost surely false, but in practice this simplification works remarkably well!

From this assumption, it follows that

$$P(S) = \frac{p_1 p_2 \cdots p_n}{p_1 p_2 \cdots p_n + (1 - p_1)(1 - p_2) \cdots (1 - p_n)},$$

where $p_i = P(S \mid W_i)$, the probability that a message is spam given that it contains W_i, and $W_1, W_2, \ldots, W_n$ represent the unique words in the email.

Due to floating-point underflow, instead of calculating $P(S)$ directly, it is better to calculate

$$\log\left(\frac{1}{P(S)} - 1\right) = \sum_{i=1}^{N}\left(\log(1 - p_i) - \log p_i\right)$$

because the summation doesn't have problems with underflow due to multiplying many small numbers together, and to then calculate $P(S)$ after a numeric expression for the right side of the above equation has been obtained.

Writing a naive Bayes spam classifier

- In R, write a naive Bayes spam classifier using the CSDMC2010 SPAM corpus training data. The `.eml` files can just be read in as plaintext files. Use the entirety of each message, including the HTML tags and the email headers.

    - For reading the text files, you may find `list.files()` and `scan()` useful.

    - You may find R's string manipulation functions to be very useful, such as `strsplit()`.

    - You can classify a message as spam if $P(S) > \frac{1}{2}$.

    - Your end goal here should be to end up with some function `detect_spam()` which can take in a string as input and decide whether or not that string is a spam email.

- Look at the words with the highest and lowest p_is. Interpret the results.

- Calculate the true positive, true negative, false positive, and false negative rates for your classifier.

- Modify your classifier so that it converts all uppercase characters to lowercase. Does this improve the performance of your classifier on the training data?

- Find examples of both spam and non-spam emails from your personal email accounts. See if your classifer classifies them correctly.

Using n-grams with logistic regression

We can compare our naive Bayes classifier with logistic regression. For our logistic regression, we will use as features the frequency counts of individual words, *i.e.*, the number of time each word we know appears in an email. Additionally, we will also use n-grams, which are sequences of n consecutive words.

- Use the `ngram` package to create a dataframe of 1-grams and 2-grams from the training data with the `ngram()` and `get.phrasetable()` functions. Each row should represent a particular email and each column should be one of the 1-grams or 2-grams.

- Use regularized elastic net logistic regression to predict spam vs. not-spam, selecting the hyperparameters α and λ with the `caret` package.

- Compute the true positive, true negative, false positive, and false negative rates for you logistic regression spam classifier. Compare its performance to that of your naive Bayes spam classifier.

Analysis of Github commit logs

We will use the Github API to scrape the commit logs for Linus Torvalds, creator of the Linux kernel, and Bram Moolenaar, creator, maintainer, and benevolent dictator for life of the Vim text editor. Afterward, we will perform sentiment analysis on their commit messages.

Using the Github API

Since Linus and Bram make most of their commits to Linux and Vim respectively, we can use the API to (1) get some of the latest commits to Linux and Vim and (2) strip out all of the commits which don't come from either of them.

The Github API can be accessed directly via your browser. In general, you begin with the url `https://api.github.com/` and then successively append text to it, *e.g.*, `https://api.github.com/users/JonahSinick`.

- Referencing the API documentation on commits, figure out which API queries will return the latest commits for Linux and for Vim. (A parameter beginning with a colon (`:`) is a *variable* which you should fill in with the appropriate value.)

- Write a Python script to access the Github API and download our desired data. Follow these specifications:

  - Use `urllib.request` to download the results of API calls. Use the `json` module, particularly the `loads()` function, to strip out all the commit messages which don't come from Linus or Brad.

  - Write the commit messages to two files, `linus.txt` and `brad.txt`, with one message per line.

4

Performing sentiment analysis

- Load the files containing the commit messages for Linus and Brad. Process them, creating one vector for Linus's messages and another vector for Brad's messages.

- Install and load the `qdap` package, which has functions for both cleaning text and performing sentiment analysis.

- Following the Cleaning Text & Debugging vignette, use `qdap` to clean the commit messages in preparation for sentiment analysis. In particular, `check_text()` should suggest the usage of some text-cleaning functions to use.

- Combine all of the commit messages into a character vector with many entries. In addition, create a vector of labels (integers 0 or 1) which indicate whether the corresponding entry in the aforementioned character vector is from Linus or Brad.

- Use `polarity()` with its default settings to perform sentiment analysis, passing in both the character vector of every commit message as well as the grouping vector.

The results of the analysis are stored in `$all`, a data frame with a column `polarity` for the sentiment polarity score of each message.

- Plot two histograms of the polarity scores for Linus and Brad overlaid on top of each other. Interpret the results.

- Look at the commit messages which had the lowest and highest polarity scores.

Latent Dirichlet allocation on Wikipedia articles

The technique of Latent Dirichlet allocation (LDA) is analogous to performing factor analysis on text. Intuitively, it reads a large collection of documents – a *corpus* – and tries to find what the "topics" of the documents are.

In Python, you will scrape every article on Wikipedia which falls into the machine learning category and process the text. Afterward, you will run LDA on the corpus of text using R.

Overview of LDA

Wikipedia gives a good intuitive explanation of LDA:

In LDA, each document may be viewed as a mixture of various topics. [...]

For example, an LDA model might have topics that can be classified as **CAT_related** and **DOG_related**. A topic has probabilities of generating various words, such as *milk*, *meow*, and *kitten*, which can be classified and interpreted by the viewer as "CAT_related". Naturally, the word cat itself will have high probability given this topic. The **DOG_related topic** likewise has probabilities of generating each word: *puppy*, *bark*, and *bone* might have high probability. Words without special relevance, such as the (see function word), will have roughly even probability between classes (or can be placed into a separate category). A topic is not strongly defined, neither semantically nor epistemologically. It is identified on the basis of supervised labeling and (manual) pruning on the basis of their likelihood of co-occurrence. A lexical word may occur in several topics with a different probability, however, with a different typical set of neighboring words in each topic.

Precisely, we first pick a number of topics k. Next, we posit a *generative model* according to the following:

1. Topics are probability distributions over the set of words used in all documents. That is, each topic is represented by the assignment of a number between 0 and 1 to each distinct word such that the sum of all those numbers is 1. For example, if we have a very simple corpus that only has the words "a" and "b", then a possible topic T would be represented by $T("a") = 0.3$ and $T("b") = 0.7$.
2. Similarly, each document is a probability distribution over the set of topics.
3. Each document has a set number of words. Each word is generated as follows: First, we randomly pick a *topic* based on the proportions of the topics associated with its document. Next, we look at the probabilities associated with that topic and accordingly randomly choose one of the dictionary words.

The algorithm optimizes the probabilities associated with topics, documents, and words so as to maximize the *likelihood* associated with the training data. For each document, the associated distribution of topics is assumed to have a Dirichlet prior, hence the name of latent *Dirichlet* allocation. (The *latent* comes from the generative model falling into the class of *latent variable models*.)

Scraping Wikipedia pages

We will first scrape all of the text which we need.

In the following, I break down the tasks which you will perform in Python into a natural series of functions and loops. You should strive to test each of

the functions you write as you write them, ensuring that the output each one returns is what you expect.

- Write a Python script to find the URL to every Wikipedia page in the machine learning category. Follow these specifications:

    - Use `urllib.request` to write a function `download_page(url)` which downloads the HTML of the page at `url` and returns it.

    - Using `download_page()`, download the Wikipedia page Category:Machine_learning. Write a function `get_urls(html)` which takes in the HTML of a Wikipedia category page and returns a dictionary with two entries: `pages`, a list of URLs to articles listed in the category page, and `subcategories`, a list of URLs to subcategories listed in the category page. Use Beautiful Soup to parse raw HTML. You can test your function on Category:Machine_learning to verify that it works.

    - Using `get_html()`, get a list of the links on Category:Machine_learning, and then add to the list the links on the *subcategories* of Category:Machine_learning, and then add to the list the links on the subcategories of the subcategories of Category:Machine_learning, and so on and so forth until there are no more subcategories to traverse.

    - Write the list of article URLs to a text file, `wp_ml_urls.txt`, with one URL per line.

- Write a Python script to download each of the Wikipedia articles in the machine learning category. Follow these specifications:

    - We need to create a folder in which we can store our downloaded HTML files. Using the `os` and `shutil` modules, check if a folder called `raw_text` exists and delete it, along with all of its contents, if it does, and then create a new folder called `raw_text`.

    - Copy and paste your `download_page()` function from before into your current script. Write a function `download_article(url)` which downloads the Wikipedia article at `url`, parses the HTML with Beautiful Soup, and returns the text of the article without any of the HTML tags. (You may find the `.get_text()` function, available for Beautiful Soup HTML objects, helpful.)

    - Read in the text file of URLs which you created earlier. Iterate over the URLs and call `download_page()` on each one. For each url `[..]/en/Page_name`, save the associated text to `raw_text/page_name.txt`; that is, take everything in the URL after the last forward slash (`/`), remove all non-alphanumeric characters, make the string completely lowercase, and use that as the file name for the article's text.

We want to remove stop words, like "the", "is", "at", and "on" from our documents. Since they show up very frequently, they will dominate the topics in LDA and our results will be useless.

Moreover, we want to group together similar words, like "apples" with "apple" or "abaci" with "abacus". We have two methods available to us: *stemming* and *lemmatization*.

Stemming will strip away everything aside from the *stem* of a word. Sometimes, the stem itself is a word, like with "cats" → "cat". However, this is not guaranteed, like with "argue", "argues", "arguing" → "argu". It is often more convenient to use lemmatization, which does not simply cut off the end of words but rather reduces different inflections into the same base form. With lemmatization, "argue", "argues", and "arguing" all map to "argue". Since the base forms are all regular English words, the results of natural processing with lemmatized words is easy to read and interpret.

To that end:

- Write a Python script to further process the text of each of the downloaded Wikipedia. Follow these specifications:

    - Write a function `process_text(text)` which reads in a string `text` and returns its processed form, where you've removed all punctuation, converted everything to lowercase, removed stop words, and lemmatized the remaining words. Use the `nltk` package to do so, referring as necessary to the documentation on downloading NLTK corpora (which includes a stop word corpus) and on `nltk.stem`. (You can use the WordNet Lemmatizer implemented in NLTK.)

    - Like earlier, check for the existence of a folder called `processed_text`. If it exists, delete it along with its contents, and then create it.

    - Use `glob` to get a list of every `.txt` file saved in the `raw_text` folder.

    - Iterate over the paths to the text files. For each one, open the file, run its contents through `process_text()`, and save the results into the `processed_text` folder with the same filename (*e.g.*, `raw_text/alphago.txt` → `processed_text/alphago.txt`).

    - Examine your processed text files to ensure that the output is what you expect it to be.

Running LDA on Wikipedia pages

We are finally (almost) ready to run latent Dirichlet allocation on our text files!

First, we need to load them into R with the `tm` package, designed for natural language processing tasks.

- Load the corpus of text into R by running

  ```
  corpus = VCorpus(DirSource(directory=d, encoding="UTF-8"))
  ```

  where d is a string containing the path to your `processed_text` directory.

- Create a document-term matrix from your corpus by running `DocumentTermMatrix(corpus)`.

- Install and load the `topicmodels` package, which has an implementation of LDA.

- Select some arbitrary number of topics to use below 50 and run LDA on your document-term matrix with `LDA(dtm, k=num_topics)`. (This can take a couple minutes.)

- The easiest way to look at the results of LDA is to see which words have the highest "loadings" on each topic. View the top n terms of each topic in your LDA model with `terms(lda_model, n)`. Interpret the results.

- Calling `topics(lda_model, n)` will show the top n topics for each document (in terms of their "loadings"). Pick a couple of article which you find interesting and look at the topics associated with those documents. Are the results surprising or expected?

- Try running LDA again with a significantly different number of topics (2x difference or greater). How different are the topics? Are they more or less coherent than before?

Finding the right number of topics is difficult and there are many different ways to do so. Each model is associated with a *log-likelihood*, so the easiest way to choose the number of topics is to do a grid search over k and choose the one associated with the highest log-likelihood. One can also consider the perplexity metric, which is essentially measuring the same thing as the log-likelihood. There is also the `ldatuning` package, which calculates four different metrics for the "quality" of a LDA model. Finally, one can use the harmonic mean method to select the optimal k. Finally, Wallace *et al.* (2009) gives some even more complex (but better) methods for evaluating the quality of topic models.

For simplicity, we'll stick to the log-likelihood.

- Choose a reasonable grid for a grid search over values of k, storing each LDA model as you go. At the end, look at their associated log-likelihoods with the `logLik()` function. Plot the log-likelihoods against the number of topics k. If you have the time, do a second, finer grid search to get a better estimate of the optimal value of k. Examine the best model you've found in comparison with the other, less optimal models, using both `terms()` and `topics()`.

In reality, what you *should* be doing is estimating the log-likelihood via cross-validation, but it would simply be far too laborious to do so now. Good to know for the future, though.

Finally, note that it often produces better results to use *n*-grams, discarding the bag-of-words assumption. Using *n*-grams is computationally challenging, but there are ways to detect *phrases* in a text preprocessing step.

- Read David Mimno's Using phrases in Mallet topic models.

Closing notes

LDA is most useful for learning structure for corpora which are too large for humans to immediately fully understand. It has many extensions, such as *correlated topic models* which allow for greater correlations between topics or *dynamic topic models* which track the evolution of topics over time, with a huge amount of the work in this field being done by David Blei.

Some interesting applications of LDA include:

- Mimno and McCallum (2007), Organizing the OCA: Learning Faceted Subjects from a Library of Digital Books

- Hu and Saul (2003), A Probabilistic Topic Model for Unsupervised Learning of Musical Key-Profiles

- Pritchard *et al.* (2000), Inference of Population Structure Using Multilocus Genotype Data, which was written before the development of LDA as it is now but proposes essentially the same generative model

Extra: Sentiment analysis on Donald Trump

If you have some free time, work through Sentiment Analysis on Donald Trump using R and Tableau, typing out all the code yourself and ensuring that you understand every line. (This is strictly optional.)