

Attributes, Factors, and Matrices

In this slightly technical lesson, we'll be discussing **attributes**.

- First, we'll begin with a high-level overview of what attributes are.
- Afterward, we'll look at two very common use cases of attributes: *factors* and *matrices*.

Attributes

Every object in R can be associated with its **attributes**, where each attribute of an object has a unique name and takes on some value.

Exercise. Use `attributes()` to view the attributes of the built-in variable `mtcars`. Try modifying some of the attributes through direct assignment – does this work? When does it fail?

Exercise. Using `attr()`, write a function that “doubles” all the names of a named list (e.g., `"colname" → "colnamecolname"`).

In general, attributes are not preserved upon object modification. However, *names*, *dimensions*, and *class* usually persist. These three attributes, being very important, have dedicated functions to access them (`names()`, `dim()`, and `class()`), which should always be used instead of `attr()`.

Now you know where names are stored!

Factors

Factors are used to represent data that can fall into one of a finite number of categories. Examples of data most naturally encoded as factors include: gender, marital status, race, nationality, and profession.

The *factor* class is built on top of integer vectors. This is the crucial insight about factors, similar to the insight about data frames being built on top of lists, and if you remember a single thing from this section, let it be that.

Factors differ from simple integer vectors in one important way: each factor is associated with a “levels” attribute, accessible through `levels()`. Each entry in a factor *must* be equal to one of its levels.

Exercise. Create a factor with `factor()` applied to an arbitrary vector, examine its levels, and try to assign a value to the factor that isn't one of its levels. Add that value to its levels (using `levels()`) and retry the assignment.

Exercise. Can you combine two factors with `c()`? If not, what's a reason why?

Hadley Wickham has the following to say about where factors often show up:

Sometimes when a data frame is read directly from a file, a column you'd thought would produce a numeric vector instead produces a factor. This is caused by a non-numeric value in the column, often a missing value encoded in a special way like `.` or `-`. **To remedy the situation, coerce the vector from a factor to a character vector, and then from a character to a double vector.** (Be sure to check for missing values after this process.) Of course, a much better plan is to discover what caused the problem in the first place and fix that; using the `na.strings` argument to `read.csv()` is often a good place to start.

Unfortunately, most data loading functions in R automatically convert character vectors to factors. This is suboptimal, because there's no way for those functions to know the set of all possible levels or their optimal order. Instead, use the argument `stringsAsFactors = FALSE` to suppress this behaviour, and then manually convert character vectors to factors using your knowledge of the data.¹

While factors look (and often behave) like character vectors, they are actually integers. Be careful when treating them like strings. Some string methods (like `gsub()` and `grepl()`) will coerce factors to strings, while others (like `nchar()`) will throw an error, and still others (like `c()`) will use the underlying integer values. For this reason, **it's usually best to explicitly convert factors to character vectors if you need string-like behaviour.**

Be sure to pay attention to these details when you load data from external sources.

Advanced R, 2.2.2.3. Define `f1 = factor(letters)`, `f2 = rev(factor(letters))`, and `f3 = factor(letters, levels = rev(letters))`. How do these three factors differ?

Exercise. Suppose you're reading in data from a file that's read into a character file into a variable `fruits`, which ends up being equal to `c("apple", "grapefruit", "NA", "apple", "apple", "-", "grapefruit", "durian")`. Read the `factor()` documentation and figure out how to convert `fruits` into a factor such that the character-encoded missing values, `"NA"` and `"-"`, end up as actual NAs. Use the `table()` function to view the `fruits` factor and verify that the `"NA"` and `"-"` strings were correctly converted into NAs.

Exercise. Write a function that takes in an arbitrary *character vector* and returns that vector as a factor, with `NA` included in the levels of the factor if it exists in the original factor.

¹Wickham also writes: "A global option, `options(stringsAsFactors = FALSE)`, is available to control this behaviour, but I don't recommend using it. Changing a global option may have unexpected consequences when combined with other code (either from packages, or code that you're source()'ing), and global options make code harder to understand because they increase the number of lines you need to read to understand how a single line of code will behave."

Exercise. Write a function that takes in a data frame `df` and converts the first `floor(ncol(df)/2)` columns into factors.

Exercise. Write a function that takes in a data frame and converts every column with at most 5 unique values into a factor. (You may find `unique()` useful.)

Exercise. Write a function that takes in a data frame and, for each factor column, replaces every NA with the most common non-NA value in the column. Generate a toy dataframe to use to demonstrate that your function works. Write a different function that replaces every NA value with a random level of the factor, distributed identically to their relative frequencies in the column's non-NA values. How can you make this [imputation](#) method reproducible? (*Hint:* Try `set.seed()`.)

Exercise. Write a function that takes in a data frame, with some but not all columns being factors, and expands each factor into a set of *indicator variables* within the data frame. Precisely, for each factor, *replace* that factor column with a number of *binary indicator variables*, having these properties:

- Every level of the factor, aside from the first level, corresponds to a new binary indicator variable.²
- For a particular row, each binary indicator variable takes on the value 1 if the original factor was equal to its corresponding level and 0 otherwise.
- For each binary indicator variable, its name is equal to the following strings concatenated together in order: (1) the name of the original factor, (2) an underscore (""), and (3) the name of the factor level itself.

You can assume that there are no NAs in the input dataframe. Test your code on the data frame `df = mtcars[1:10,]; for (n in c("cyl", "am", "carb")) df[[n]] = factor(df[[n]]);` you should obtain a result with 13 columns.

To illustrate the desired functionality, consider the following diagram:

col		col_2	col_3
---		-----	-----
1		0	0
2		1	0
1	=>	0	0
2		1	0
2		1	0
3		0	1

²The reason for not making a binary indicator variable for the first level of the factor is that the state of "factor is equal to its first level" can already be represented by the `length(levels()) - 1` binary indicator variables all being set to 0. As such, if we add in a binary indicator variable for the first level, it will cause problems with collinearity that breaks linear techniques (including regression).

The motivation behind this functionality is that linear regressions can't be directly run with factors, which are *categorical* variables, in the predictors; the integer-valued `levels()` of a factor don't encode any meaning. Instead, we separate all but one level out into *binary indicator variables*, which we can regress against. The simplest and most common example of this is encoding a gender variable as 0 or 1.

Exercise. Use `load()` to load `time.dat`, a two-column subset of the data from the [National Longitudinal Study of Adolescent Health](#). (The function will load it into the variable `df`.) Look at the documentation for [Wave II: In-Home Questionnaire, Public Use Sample](#) and read about the two questions in the data (check the column names). Write some code to convert each column to numeric values representing “number of hours past 8:00 PM”³ and plot two histograms, one for each column, overlaid on top of each other using multiple `geom_histogram()` calls with the `fill` and `alpha` parameters. (*Hint:* Move the `aes()` call into the `geom_histogram()` calls.)⁴

Matrices

Arrays are simply atomic vectors with a *dimension attribute*, which must be a vector of integers. (The size of the vector must be compatible with its dimensions.) **A matrix is an array with two dimensions.** Matrices can be created as such:

```
> matrix(1:100, nrow=10)
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
[1,]    1   11   21   31   41   51   61   71   81   91
[2,]    2   12   22   32   42   52   62   72   82   92
[3,]    3   13   23   33   43   53   63   73   83   93
[4,]    4   14   24   34   44   54   64   74   84   94
[5,]    5   15   25   35   45   55   65   75   85   95
[6,]    6   16   26   36   46   56   66   76   86   96
[7,]    7   17   27   37   47   57   67   77   87   97
[8,]    8   18   28   38   48   58   68   78   88   98
[9,]    9   19   29   39   49   59   69   79   89   99
[10,]   10   20   30   40   50   60   70   80   90  100
```

Exercise. Write a function that takes an n -by- m numeric matrix and turns it into a m -by- n numeric matrix where both matrices give the same result when `as.numeric()` is applied.⁵

³Deciding what the cutoff point is where you stop counting something as “X hours after 8 PM” and start counting it as “24-X hours before 8 PM” is an arbitrary choice in principle, but some choices are more sensible than others – use your human judgment to optimize for interpretability!

⁴You can use string manipulation or R's [inbuilt time classes](#). Try doing it both ways. For plotting the overlaid histograms, try `ggplot(df) + geom_histogram(aes(x=H2GH42), fill="red", alpha=0.2) + geom_histogram(aes(x=H2GH43), fill="blue", alpha=0.2)`.

⁵Just modify the attributes.

Matrices work similarly to data frames: you have access to `colnames()`, `rownames()`, `ncol()`, and `nrow()`, and subsetting (broadly speaking) works the same way.⁶

Exercise. Run the following code: `df = data.frame(matrix(1:100, nrow=10)); df[5, 5] = NA; df[6, 6] = NA`. Figure out how `df[is.na(df)]` works. Write and test a function that takes as input a data frame `df` of purely numeric data and a number `k`, returning a vector of every number in `df` divisible by `k`.

You won't work directly with matrices directly very much for most data science applications – data frames are much more common and important.⁷ In practice, they mostly show up as *intermediate forms* in the conversion and manipulation of data, and it's helpful to be aware of matrices so you can debug problems when they show up. (Knowing how matrices work in R might be more important if you're doing numerical simulations of some sort, perhaps in the computational science.)

Some more supplementary exercises on matrices:

Exercise. If you aren't familiar with the operation of *matrix multiplication*, read about it briefly on [Wikipedia](#). Does matrix multiplication work normally with the `*` operator? Why or why not? (Try testing multiplication with the identity matrix.)⁸

Exercise. Write a function `min_matrix(n, m)` with `n` rows and `m` columns where the value in row `i`, column `j` is equal to `min(i, j)`.

Exercise. Write a function to determine if an input matrix is symmetric across its main diagonal (so the element in row `i`, column `j` is equal to the element in row `j`, column `i`). Only square matrices can be symmetric; you may find the `t()` function helpful.

Exercise. Write a function `trace(mat)` to calculate the trace of an input matrix, *i.e.*, the sum of its diagonal elements. (You may find `diag()` helpful.) Check whether or not the `trace()` function is *linear*, that is, whether it's true that `trace(A + B) = trace(A) + trace(B)` and `trace(c*A) = c*trace(A)` for some constant `c`.

Exercise. Consider the matrices given by `mystery = function(x) matrix(c(cos(x), -sin(x), sin(x), cos(x)), nrow=2)`. Is the `mystery()` function *periodic* in some sense? Check if this holds in practice; if there's a

⁶There are some nuances – if you want to learn about them, read Hadley Wickham's *Advanced R*.

⁷Isn't it odd that the majority of expositions of R focus so much on presenting matrices in the very beginning, even though they are far more removed from doing *actual, interesting work* than data frames and even many other R concepts?

⁸Matrices are vectors, so the multiplication is done element-by-element. Proper matrix multiplication is done with the `%*%` operator. It's not immediately useful to know this, but it's good to keep it in mind because for more complex matrices, it may not be immediately obvious that `*` doesn't correspond to ordinary matrix multiplication.

discrepancy, explain it. (*Hint*: If you aren't familiar with the behavior of the `sin()` and `cos()` functions, graph a scatterplot of their values against `seq(0, 10, 0.01)`.)⁹

Exercise. Write a function to turn lists of length 4 into 2-by-2 matrices, forming a list-matrix capable of holding different data types. (*Hint*: Use `dim()`.) Speculate on some use cases of list-matrices.

Optional challenge exercises

Feel free to skip these – they're optional.

Exercise. Write a function `matrix_mult(A, B)` that implements matrix multiplication, computing `A %*% B` (without using the `%*%` operator, of course). Compare the performance of `matrix_mult()` with the built-in `%*%` operator with the `timeit` function for matrices of different sizes. Can you speed up your code?

Exercise. On the 2D plane, we can identify the *point* (x, y) with the *column vector* `matrix(c(x, y), nrow=2, ncol=1)`. First, write a function which takes in a list of column vectors, internally adds each one as a row of a two-column dataframe, and then uses `ggplot()` to graph a *scatterplot* of all the points in the input list. Second, modify this function to accept an argument `x` and so that for every column vector `c` in its list of points, instead of putting the data in `c` directly into the data frame, it does so for the product `mystery(x) %*% c`. Experiment with different values of `x` and come up with a geometric understanding of how the output graph changes as you modify `x`.¹⁰

Exercise. Construct several arbitrary square matrices of identical size but differing contents (try to make them interesting, with lots of nonzero numbers), and consider the trace of their *product*. If you change the order in which they're multiplied, does the trace necessarily stay the same? Characterize the permutations of the order of matrices in the product which leave the trace of the product unchanged. (It may be instructive to consider all permutations of 3 different matrices.) What about the case where the matrices are symmetric?

⁹We have `mystery(0) != mystery(2*pi)` because of floating-point imprecision.

¹⁰Calling `mystery(x)` returns a 2-dimensional rotation matrix corresponding to rotation through the angle `x` (given in radians).