Functional Programming in R

> And he to me: "Remember now your science, which says that when a thing has more perfection, so much the greater is its pain or pleasure."" — Dante, *Inferno*, 6. 106-109

So far, you've been using *for* and *while* loops in R for iteration. There are, however, benefits to a functional programming approach.

In an iterative style, we *loop* through values and successively manipulate each value, whereas in a functional style we *apply* some function to every value independently. It's easiest to illustrate with an example.

Suppose that I have the following dataframe:

```
> df = data.frame(matrix(1:100, nrow=10))
> df
   X1 X2 X3 X4 X5 X6 X7 X8 X9 X10
1   1 11 21 31 41 51 61 71 81  91
2   2 12 22 32 42 52 62 72 82  92
3   3 13 23 33 43 53 63 73 83  93
4   4 14 24 34 44 54 64 74 84  94
5   5 15 25 35 45 55 65 75 85  95
6   6 16 26 36 46 56 66 76 86  96
7   7 17 27 37 47 57 67 77 87  97
8   8 18 28 38 48 58 68 78 88  98
9   9 19 29 39 49 59 69 79 89  99
10 10 20 30 40 50 60 70 80 90 100
```

Now, perhaps I would like to calculate the mean of every column.

One way to do this is to loop through the columns and use the `mean()` function:

```
> means = c()
> for (i in 1:ncol(df)) {
+    means = c(means, mean(df[[i]]))
+ }
> means
 [1]  5.5 15.5 25.5 35.5 45.5 55.5 65.5 75.5 85.5 95.5
```

However, I can do this in a somewhat more compact fashion by using R's `sapply()` function:

```
> means = sapply(1:ncol(df), function(i) mean(df[[i]]))
> means
 [1]  5.5 15.5 25.5 35.5 45.5 55.5 65.5 75.5 85.5 95.5
```

In general, the family of `*apply()` functions in R all facilitate programming in a functional paradigm.

Using `sapply()` with anonymous functions

We'll first learn about functional programming by using `lapply()`. The other `*apply()` functions are mainly extensions of `lapply()`, and we'll cover them later.

A picture is worth a thousand words:

Figure 1: A visual illustration of `lapply()`.

Here's an example of using `lapply()` to double every number in a vector. Run the following code:

```
double = function(x) {
  2*x
}
lapply(1:10, double)
```

We first create a function `double(x)` which returns 2 times its argument, and then we `lapply()` the `double()` function onto the vector `1:10`, with the result of each computation returned in a list. In general, when calling `lapply(values,`

2

func), each value of `values` is supplied as an unnamed first argument to func().

Exercise. Why might we want to by default return the output of `lapply()` in a *list* instead of just `unlist()`ing the values automatically?[1]

We can write this more compactly using an *anonymous function*, which is an unnamed function defined for use in a local context only. Run the following code:

```
> lapply(1:10, function(x) 2*x)
 [1]  2  4  6  8 10 12 14 16 18 20
```

If we anticipate that we won't be using a function often enough to give it a name, we can define it within `sapply()` like we did with `function(x) 2*x`. (Recall that a function doesn't need an explicit `return()` statement – it returns the last expression evaluated by default – and that I only need curly braces for a multiline function definition.)

Exercise. Write a function using `lapply()` and `class()` to print out the class of each column in the built-in `mtcars` dataset. Run `unlist()` at the end so it prints in a more human-readable format. (*Hint:* Remember that data frames are built on top of lists.)

Exercise. Write a function using `lapply()` to standardize each column of `mtcars` by (1) subtracting off its mean and (2) dividing it by its standard deviation (given by `sd()`). Be sure to check that your function returns a data frame.

Exercise. Write a function using `lapply()` that standardizes every numeric column of an input data frame and leaves the others unchanged. Test your function on the dataframe defined by `df = data.frame(matrix(1:100, nrow=10))`; `df[1:5] = lapply(df[1:5], as.character)` (understand what this code is doing as well).

Looping patterns

We'll pause for a moment to discuss, at a higher level, the operation of looping.

In general, there are three main ways to loop through a list-based data structure:

- Looping through the elements: `for (col in df)`

- Looping through the indices: `for (i in 1:length(df))`

- Looping through the names: `for (n in names(df))`

---

[1]Functions in R don't have a return type, so we don't know in advance what they'll return. Although the `double()` function only returns numerics, that isn't always the case, so it's best to return results in a `list()`, which allows for multiple types in its entries.

Exercise. The `lapply()` equivalent of the first is `lapply(df, function(x) {})`. Write down equivalents for the second and third forms.

The first form is the simplest, but you don't get the name or index of each item. The second and the third are more complex, but provide you with more information, so keep them in mind – they may be helpful for more complex problems.

`vapply()` and `sapply()`

`sapply()` is in fact not the most basic of the `*apply()` functions; `lapply()` is. Here's a brief description of two more functions to give you a sense for the overall landscape:

- `lapply()` maps a function onto a list and *returns a list*.

- `vapply()` is an extension of `lapply()` that maps a function onto a list and *returns an atomic vector*. It takes an additional argument specifying the *type* and *length* of each element of the return vector, throwing an error if they don't match.

- `sapply()` is an extension of `lapply()` which will `unlist()` the results. If appropriate, it will also assign dimensions to the output, turning it into a matrix.

Let's look at `vapply()`. It's called as such (run this code):

```
vapply(mtcars, class, character(1))
vapply(list(matrix(1:9, nrow=3), matrix(1:20, nrow=5)), dim, numeric(2))
```

In general, when calling `vapply(args, func, example)`, *each time func() is called on an element of args*, the output must have the same type and length as `example`. Otherwise, `vapply()` stops with an error.

Exercise. Using any inputs you like, write valid `vapply(args, func, example)` calls which have each of the following values for `example`: `logical(3)`, `numeric(10)`, `character(2)`.

Exercise. With a variety of different functions, test the behavior of `vapply()` and `sapply()` when the list of arguments is an empty list (`list()`). How would the behavior of `vapply()` help you write code robust to errors and bugs?[2]

Exercise. What happens when `sapply(args, func)` is called in a situation where `func()` returns vectors of different lengths for different elements of `args`? How can `vapply()` be used to detect unexpected instances of this situation?

Exercise. Experiment with lists containing `Sys.time()`.

---

[2]Since `vapply()` will return the correct type of 0-length vector in the case where the list of arguments is empty, it helps guard against errors from various edge cases.

It's dangerous to use `sapply()` when writing functions you'll use elsewhere, because you won't know if your output is an unexpected type or has an unexpected length until your program exhibits strange behavior elsewhere. It's better to use `vapply()`, which throws an error when the output isn't of the specified type and length and enforces type consistency in various edge cases. However, it's fine to use `sapply()` when working interactively in the console, where you'll be able to visually notice any strange behavior.

Passing in named arguments

If you have `sapply(df, func)` and want to pass in named arguments to every call of `func()`, you can do so by passing in named arguments into `sapply()` directly, *e.g.*, `sapply(df, func, param=TRUE)` will call `func(c, param=TRUE)` for every column `c` of `df`.

Remark. The same syntax works for `lapply()`. For `vapply()`, the named arguments go after the example return value.

Exercise. Write a function using `sapply()` to find the mean of every vector in a list of numeric vectors, ignoring `NA` values. Test your function on the list `L = lapply(1:5, function(x) sample(c(1:4, NA)))`.

Exercise. Let `trims = seq(0, 0.5, 0.1)` and `x = rnorm(100)`. Rewrite the expression `lapply(trims, function(trim) mean(x, trim=trim))` to not need an anonymous function.

Why use `*apply()` instead of loops?

At times, the usage of loops is inevitable and the most natural way to program something. Don't get caught up in trying to code something functionally if a loop seems intuitive. In particular, these three use cases are more suitable for loops than for functional programming:

- Modifying a data structure in place (changing it without making a new copy)
    - This is because you have to use the `<<-` operator to modify the object while situated in the scope of a function call.
- Recursive functionality
- While loops

Moreover, people will sometimes say say that you should use the `*apply()` functions instead of loops because loops are slow. **This is not true.**

As we saw earlier with the n-dominoes problem, loops can be sped up significantly by *preallocating memory* for the data structures which you access. In

general, loops can be made approximately as fast as writing equivalent code for a function to be used with *apply() if you follow these guidelines[3]:

- Initialize new objects to full length before the loop, rather than increasing their size within the loop.

  – Every time you increase the size of an object within a loop, you actually *copy the whole structure over to a different part of memory* every single time.

- Do not do things in a loop that can be done outside the loop.

- Do not avoid loops simply for the sake of avoiding loops (see the above criteria).

Given that loops don't actually have performance issues in R, why should we use *apply() functions at all? For these three reasons:[4]

- Using the *apply() functions can make it clearer what you're doing.

  – The notion of applying the same function to every element of a list is in general very intuitive. Code clarity is important, both for yourself and for others.

- The *apply() functions have no unwanted side effects.

  – That is to say, their functionality is *isolated* from the rest of your code, so it's harder for you to make accidental modifications to variables you've defined elsewhere.

  – Two caveats are that this isn't true if you use assign() or the <<-operator, which are seldom used and only show up in very specific situations.

- Multi-core processing packages for R implement parallelization by over-writing the built-in *apply() functions with their own versions. As such, liberal usage of *apply() in your code means that you'll be able to easily parallelize it without much rewriting.

  – It's precisely because the *apply() functions have no side effects that they're often used for parallelization purposes. Otherwise, it can be very difficult to ensure correct program behavior.

  – More importantly, when calling, say, sapply(args, func), each call of func() is completely independent of the other ones. This allows them to easily be dispatched to different processor cores.

---

[3]From a 2008 issue of R News.

[4]See the answers to Is R's apply family more than syntactic sugar, including the comments on the first one.

Improving old code

Exercise. Go back to your old code for various R exercises. Find five functions which could be written more easily or clearly using functional programming (*e.g.,* using the *apply() functions instead of a for loop) and rewrite them. Check the difference in runtime.

Exercise. If you remember writing any R code in the past for exercises where you kept adding values to a vector or list on every iteration of a loop, rewrite the code here using a preallocated data structure (if possible). CHeck the difference in runtime using the timeit package.

Supplemental exercises