Regularized Linear Regression

Some helpful notes on the `glmnet` package are at the end of this document. Look there if you have problems.

Exploring regularization with simulated data

Define x and y using:

```
set.seed(1); j = 50; a = 0.25
x = rnorm(j)
error = sqrt(1 - a^2)*rnorm(j)
y = a*x + error
```

If you run `summary(lm(y ~ x - 1))`, corresponding to a linear model with no constant coefficient, you should get an estimated value of 0.2231 for a.

Write a function `cost(x, y, aEst, lambda, p)` which takes

- Two vectors x and y of equal length
- An estimate of the value of a, aEst
- A regularization parameter lambda
- A number p = 1 or 2, indicating whether L^1 or L^2 regularization is being performed

and returns the L^p regularized cost function associated with the estimate y = aEst*x.

Create a dataframe with two columns, one corresponding to the values of λ $2^{-8}, 2^{-7}, \ldots, 2^{-1}, 2^0, 2^1$, and for each of these values of λ, values of a from -0.1 to 0.3, in equally spaced increments of 0.001. Use `expand.grid()` to fill in the grid.

Add `"costL1"` and `"costL2"` columns, where we'll store the cost of associated with each pair (`lambda, a`), for each of p = 1 and p = 2.

For each of p = 1 and p = 2,

- Use lapply to make a `plots` list with 10 `ggplot()` objects, one for each value of lambda from 2^{-8} to 2^{-1}, graphing values of a on the abscissa

1

(x-axis) and values of the cost function on the ordinate (y-axis). Then use multiplot with `plotlist = "plots"` to display these graphs in 2 columns of 5.

- Pay special attention to the values on the y-axis, which vary from plot to plot.

Comparing regularization and stepwise regression

We'll continue using the simplified speed dating dataset from yesterday. For now, please restrict analyzing *attractiveness ratings* (`"attr_o"`) for *males*.

Using the entire dataset

The `glmnet()` and `cv.glmnet()` functions can perform both L^1 and L^2 regularized linear regression as well as a mix of the two (which we'll be exploring later). This behavior can be tuned via the `alpha` parameter; read the official documentation to figure out it works.

- Use backward stepwise regression to generate attractiveness predictions for the whole dataset. (Don't use cross-validation at this point.)

- Use `glmnet()` to generate similar predictions with both L^1 and L^2 regularized linear regression.

    - Look at which values of λ were used by `glmnet()`.

    - Write a function that (1) takes the model object generated by a call to `glmnet()` and the true values for the target variable, (2) uses `predict()` to generate predictions for every value of λ which `glmnet()` tried, and (3) returns the λ corresponding to the lowest RMSE and the RMSE itself.

- Compare the minimum RMSE for both regularized fits with the RMSE for backward stepwise regression.

- Compare and interpret the coefficients for L^1 and L^2 regularized linear regression using the optimal values of λ determined earlier.

Making cross-validated RMSE predictions

As you saw in the assignment on resampling, we want to use *cross-validation* to get more accurate estimates of model quality. In particular, stepwise regression tends to *overfit*, because of problems with multiple hypothesis testing, so non-cross-validated estimates of a stepwise regression model's quality are often

overly optimistic. (However, it's easy to understand and, pedagogically, a good stepping stone to regularization, which is why we include it in our curriculum.)

- Briefly skim Stopping stepwise: Why stepwise and similar selection methods are bad … for an introduction to some of the problems with stepwise linear regression.

Write a function following these specifications:

- Use 10-fold cross validation to generate predictions for attractiveness with (1) stepwise regression, (2) L^1 regularized linear regression, and (3) L^2 regularized linear regression.

- For regularized linear regression, use `cv.glmnet()` to get cross-validated estimates of the optimal value of λ. As such, when generating predictions for an regularized linear model `fit`, use the value of λ stored in `fit$lambda.min`.

- Return the RMSE associated with each of the three sets of predictions.

Here are some points to keep in mind:

- Within each cross-validation fold, you'll want to `scale()` the features which you pass into `cv.glmnet()`. When generating predictions on the *held-out* data, you want to scale the features in the same way (*i.e.*, by applying the same linear transformation). The output of `scale()` will contain *attributes* which can be accessed and passed into successive calls of `scale()` to perform the same transformation.

- If you have a string, say, `"attr_o"`, and you want to pass that into `lm()` as part of the regression formula, you can paste together the formula's components (*e.g.*, `paste("attr_o", "~.")`), call `formula()` on the string to turn it into a *formula*, and then passing the formula into `lm()`.

Use your function to explore the difference in model quality between backward stepwise regression, L^1 regularized regression, and L^2 regularized regression when predicting attractiveness ratings.

Elastic net regression

Instead of penalizing the sum of squared residuals by the L^1 or L^2 norm of the regression coefficients, we can penalize with a combination of the two, corresponding to setting the `alpha` parameter in `glmnet()` to a value between 0 and 1. We can use cross-validation to find the optimal *pair* of *hyperparameters* (α, λ).

Thankfully, we won't have to implement that ourselves (for now)! Instead, we can use the `caret` package to get a cross-validated estimate of the optimal (α, λ).

Here's an example of how to use the `caret` package's `train()` function:

```
param_grid = expand.grid(.alpha = 1:10 * 0.1, .lambda = 10^seq(-3, 0, length.out=10))
control = trainControl(method="repeatedcv", number=10, repeats=3, verboseIter=TRUE)
caret_fit = train(x=features, y=target, method="glmnet", tuneGrid=param_grid, trControl=c
```

In the above example, we perform *10-fold cross-validation* repeated *3 times*.

Whereas `cv.glmnet()` picks an appropriate sequence of λ values, the `caret` package does not.

- Run `cv.glmnet()` on the data for any rating to get a rough sense for what the range of λ should be, and then use that when doing grid search with `caret`.

Write a function according to the following specifications:

- Use the `caret` package, following the above example, to find the optimal values for (α, λ) when predicting attractiveness ratings.

Finally:

- Calculate the corresponding RMSE and compare the different RMSEs for all combinations of (gender, rating).

- Now that you've gotten some experience with `glmnet()` and `cv.glmnet()`, go back and modify your cross-validation function so that you can pass in (1) any of the five rating variables and (2) a gender variable to filter by instead of restricting solely to attractiveness ratings for males.

A note on `glmnet`

Here, I'll cover two important points about the behavior of the `glmnet` package.

Passing in data

For `lm()`, you passed in the entire data frame, including both target variable and predictors. `glmnet(features, target, ...)` and `cv.glmnet(features, target, ...)` expect a *scaled matrix of predictors* for `features` and a numeric vector for `target`. The `scale()` function returns a matrix, so you can just call `scale()` on a data frame of predictors and pass that in as `features`'.

Picking values of λ

"Ordinarily", one might expect that, for every different value of λ we want to try using with regularized linear regression, we would have to recompute the entire model from scratch. However, the `glmnet` package, through which we'll be using regularized linear regression, will automatically compute the regression coefficients for *a wide range of λ values simultaneously.[1]

When you call `glmnet()` – or, later, `cv.glmnet()` – you'll get out an object, `fit`. (You should generally not be specifying *which* λ values the algorithm should use at this point – it'll try to determine that on its own.) By printing out `fit` in the console, you can see which values of λ were used by `glmnet`.

When you want to make predictions with this `fit` object, you'll have to specify *which* value of λ to use – instead of calling `predict(fit, new_data)`, you'll want to call `predict(fit, new_data, s=lambda)` for some particular λ = `lambda`. Similarly, when extracting coefficients, you'll want to call `coef(fit, s=lambda)`.

Finally, `cv.glmnet()` will use *cross-validation* to determine `fit$lambda.min` and `fit$lambda.1se`. The former is the value of λ (out of all those the algorithm evaluated) which minimizes the cross-validated mean squared error (MSE), and the latter is the greatest value of λ (again, of those evaluated by `glmnet`) such that the MSE corresponding to `fit$lambda.1se` is within 1 standard error of the MSE corresponding to `fit$lambda.min`.

If it turns out that the optimal value of λ lies at either end of the range of λ values used by `glmnet`, then you'll want to modify the range of λ. However, the documentation advises against passing in just a single value for the `lambda` parameter of `glmnet()` and `cv.glmnet()`, instead suggesting modifying `nlambda` and `lambda.min.ratio`.[2] Nevertheless, there are times when passing in a single value makes sense, like when you've previously determined the optimal λ and want to just use that instead of a range of different λ values.

---

[1]"The `glmnet` algorithms use cyclical coordinate descent, which successively optimizes the objective function over each parameter with others fixed, and cycles repeatedly until convergence. The package also makes use of the strong rules for efficient restriction of the active set. Due to highly efficient updates and techniques such as warm starts and active-set convergence, our algorithms can compute the solution path very fast."

[2]"Typical usage is to have the program compute its own `lambda` sequence based on `nlambda` and `lambda.min.ratio`. Supplying a value of `lambda` overrides this. WARNING: use with care. Do not supply a single value for `lambda` (for predictions after CV use `predict()` instead). Supply instead a decreasing sequence of `lambda` values. `glmnet` relies on its warms starts for speed, and it's often faster to fit a whole path than compute a single fit.""