

Self-Assessment

We'll be having another self-assessment. As before,

- Type your answers in a new R script file with comments indicating where the answer to each question begins.
- Write down the current time. Please email us (at signaldatascience@gmail.com) with your R script attached when you finish along with the amount of elapsed time.
- Work individually. You can however consult R documentation, look at old assignments, use the Internet, etc., but don't copy and paste code verbatim.
- Make your code as clear, compact, and efficient as possible. Use everything that you've learned! **Please comment your code and write it cleanly so I know what each part of your script is supposed to do.**
- If it's 12:30 and you haven't finished: send us your code and go eat lunch! It's okay if you don't finish in the allotted time.

Part 1: Regularization

Part 2: Probability

In Part 2, we'll take a break from the usual data science problems and work on some more probability!

We **won't** be going over the theoretical solutions to these questions during class time. However:

- If you have time to spare, please do think about *why* the numerical results are what they are.
- **After 7 PM**, I'll give a short (30-minute) explanation of all three probability questions to anyone who sticks around

Hashmap collision

From “120 Interview Questions”, we have the following:

Your hash function assigns each object to a number between 1:10, each with equal probability. With 10 objects, what is the probability of a hash collision? What is the expected number of hash collisions? What is the expected number of hashes that are unused?

Use `sample(..., replace=FALSE)` to give estimates of all three.

(Clarification: If n objects are assigned to the same hash, that counts as $n - 1$ hash collisions.)

Rolling the dice

Here’s a nice question from one of [the first cohort’s students](#):

Given a fair, 6-sided dice, what’s the *expected number of rolls* you have to make before each number (1, 2, ..., 6) shows up at least once?

Write code to estimate the answer.

Bobo the Amoeba

Here’s a problem commonly found in quantitative finance interviews, an easier version of which sometimes appears in data science interviews:

Bobo the amoeba can divide into 0, 1, 2, or 3 amoebas with equal probability. (Dividing into 0 means that Bobo dies.) Each of Bobo’s descendants have the same probabilities. What’s the probability that Bobo’s lineage eventually dies out?

I’ll outline a computational path for you to estimate the solution to this problem. In particular, we’re going to simulate a large number of amoeba lineages over time to determine how the probability of total extinction changes as we iterate forward in time.

- Write a function `next_gen(n)` that takes in an initial number of amoebas n , determines how many amoebas are in the next generation according to the probability above, and returns that value. Sanity check: `next_gen(1)` should return 0, 1, 2, or 3 with equal probability.
- Note the enormous computation time required for `next_gen(n)` when n is very large. If there are a *large* of amoebas, we can assume (with reasonable confidence) that the lineage isn’t going to die out. Pick a reasonably large value of n , like 500 – let’s call it N – and modify `next_gen(n)` to just return $N+1$ when $n > N$. (This is fine because we just want to know if the lineage

will *die out* or not – how huge the population can get in cases where it doesn't don't really matter to us.)

- We're going to simulate `num_lineages` lineages for `n_gens` generations, so set `num_lineages = 10000` and `n_gens = 30`.
- Initialize a matrix of appropriate size and dimensions, where each column represents a single lineage of amoebas and every row represents a different generation. Next, set the initial generation to a population of 1.
- Iterate over the number of generations. For each iteration, apply `next_gen(n)` to the population of the most recent generation to get the population for the next generation, filling in the values of your population matrix.
- Turn your population matrix into a matrix of 1s and 0s corresponding to whether the lineage was still alive or died out at each step of the simulation.
- Calculate the probabilities of lineage extinction from your population matrix. *Hint:* Take the `rowSums()` of the matrix.
- `qplot()` the time evolution of the extinction probability. What do you think it is? Give a numerical estimate using your calculations.

Now, use [WolframAlpha](#) to solve the cubic equation $p = \frac{1}{4} + \frac{1}{4}p + \frac{1}{4}p^2 + \frac{1}{4}p^3$. One of the solutions will numerically correspond to your calculated probability. Which one? Why?