

Classification

Although you have already been introduced to classification via logistic regression, we will develop the theory of classification in general in greater detail.

We will begin with some theoretical exposition and then move into a series of exercises with different classification techniques.

Overview

The task of *classification* is fundamental to the history and development of machine learning because the simplest machine learning problem possible is binary classification: determining whether data falls into one class or the other.

Generative vs. discriminative methods

The novice data scientist can easily fall into the trap of considering different techniques to be more-or-less disconnected and unrelated, as if machine learning were a “grab bag” of models and algorithms, each with its own odd little quirks and nuances. This is not the case! There is **structure** in this world.

Of particular importance is the distinction between **generative vs. discriminative** approaches. A generative approach models the way in which the observed data were generated in order to perform classification on new data, whereas a discriminative approach doesn’t care about how the data was generated: it is *purely* concerned with the problem of classifying new data.

To make this precise, suppose that we wish to predict the label y from the training data x . In a *discriminative* approach, we are concerned with evaluating $y^* = \arg \max_y P(y | x)$, that is, finding the most likely class label y^* given the data x . The only concern is how we can distinguish one class from another using x .

However, we can use Bayes’ rule to rephrase our problem. Since $P(y | x) = P(x | y)P(y)/P(x)$, we have $y^* = \arg \max_y P(x | y)P(y)/P(x)$. However, the value of y which maximizes $P(x | y)P(y)/P(x)$ is the same as the value of y which maximizes $P(x | y)P(y)$, because $P(x)$ does not change with y .

As such, we can reformulate the problem as the evaluation of $y^* = \arg \max_y P(x | y)P(y)$. This corresponds to a *generative* model, because we are directly estimating the value of $P(y)$, that is, explicitly modeling the distribution of each class in order to perform classification.

Philosophical importance

Instead of just being two equivalent formulations of the same underlying problem, the distinction between generative and discriminative modeling is very philosophically fundamental. To quote [V. N. Vapnik](#):

The philosophy of science has two different points of view on the goals and the results of scientific activities.

1. There is a group of philosophers who believe that the results of scientific discovery are the real laws that exist in nature. These philosophers are called the *realists*.
2. There is another group of philosophers who believe the laws that are discovered by scientists are just an instrument to make a good prediction. The discovered laws can be very different from the ones that exist in Nature. These philosophers are called the *instrumentalists*.

The two types of approximations defined by classical discriminant analysis (using the generative model of data) and by statistical learning theory (using the function that explains the data best) reflect the positions of realists and instrumentalists in our simple model of the philosophy of generalization, the pattern recognition model. Later we will see that the position of philosophical instrumentalism played a crucial role in the success that pattern recognition technology has achieved.

Considering classification alone, we can state this in a more precise fashion. Suppose that we have a black box, \mathcal{B} , which when given an input vector \mathbf{x}_i returns an output $y_i = -1, +1$. Given the training data (y_i, \mathbf{x}_i) , $i = 1, \dots, n$, the task of *binary classification* is to find a function which best approximates whatever rule the black box \mathcal{B} is internally using to perform the operation $\mathbf{x}_i \mapsto y_i$.

There are two different ideas of what a *good approximation* means:

1. A good approximation of \mathcal{B} is a function that is *similar in function space*, like how $\sin 2x$ is similar to $\sin 2.1x$, to the function which \mathcal{B} uses to map $\mathbf{x}_i \mapsto y_i$.
2. A good approximation of \mathcal{B} is a function that is *numerically similar* to the function which \mathcal{B} uses, giving approximately the same error rate of classification as \mathcal{B} itself.




Figure 1: The linear line is a good approximation in the first sense; the curvy polynomial is a good approximation in the second sense.

The first definition is concerned with approximating the *true function*. The second definition is only concerned with *minimizing the error rate*.

Practical importance

This distinction is not just of philosophical importance; it carries practical significance:

[I]ll-posed problems can occur when one tries to estimate *unknown reasons from observed consequences*.

That is to say, problems that are very computationally difficult to solve – where a small change in the observed data can lead to (discontinuously) large changes in the solution – are very often precisely the problems which occur when we try to use a *generative* model. If the solver possesses “very strong prior knowledge about the solution”, or if the solver makes very strong prior *assumptions*, then generative models become much more tractable, but this is not always feasible.

Vapnik writes:

Keeping in mind the structure of ill-posed problems our problem of finding the [solution for \mathcal{B}] can be split into two stages:

1. Among a given set of admissible functions find a subset of functions that provides an expected loss that is close to the minimal one.
2. Among functions that provide a small expected loss find one that is close to the \mathcal{B} function.

The first stage does not lead to an ill-posed problem, but the second stage might (if the corresponding operator is unstable).

The realist view requires solving both stages of the problem, while the instrumentalist view requires solving only the first stage and choosing for prediction any function that belongs to the set of functions obtained.

The shift from the realist view to the instrumentalist view marked a significant advance in the development of machine learning.

Evaluation of classifier performance

- Read [Nina Zumel's primer](#) on the evaluation of classifier performance.

Preliminary steps

We will mainly be using simulated data for the exercises below. First, we must write functions which generate our simulated data.

For simplicity, we will restrict attention to the 2-dimensional [unit square](#) ($x \in [0, 1], y \in [0, 1]$).

- Write a function `lin_pair(m, b, label)` that takes in integers `m`, `b`, and `label` and returns two numerics `c(x, y)` satisfying the following criteria:
 - Both `x` and `y` should be between 0 and 1.
 - If `label` is set to 1, then `y` must be greater than `m*x + b`. Conversely, if `label` is set to -1, then `y` must be less than `m*x + b`.

You can think of the function `lin_pair(m, b, label)` as picking a point in the unit square uniformly randomly from the region falling above or below the line $y = mx + b$.

- Plot the points you get from running `lin_pair()` many times with the same input parameters to visually verify that your function works correctly.

- Write a function `quad_pair(a, b, label)` which does the same thing except for the quadratic parabola $y = (x - a)^2 + b$. Verify that it works by plotting the results.
- Write a function `mvnorm_pair(mu, cov)` which returns a point sampled from a multivariate normal distribution. `mu` should be a vector containing the mean for each dimension and `cov` should be a 2-by-2 covariance matrix (where `cov[1, 1]` is the variance in dimension 1, `cov[2, 2]` is the variance in dimension 2, and `cov[1, 2] = cov[2, 1]` are the covariance of the two dimensions). You may find `mvnrm()` from the MASS package useful.
 - Note that the `select()` function from MASS will mask the `select()` function from dplyr if you load MASS after dplyr. You can access the dplyr `select()` function with `dplyr::select()`.

In the exercises to follow, we will restrict consideration almost entirely to the problem of **binary classification** for simplicity's sake. Multiclass classification is possible with many of the following methods, but it is generally wise to work with simple examples first to gain intuition and understanding.

k-Nearest Neighbors

Discriminant methods

[Ronald Fisher](#)'s method of *discriminant analysis* is one of the earliest classification methods.

The idea behind discriminant analysis is that we posit a *generative model* for our data, where each of the two classes is generated from a multivariate normal distribution, and attempt to estimate the parameters of those distributions with the training data. With the parameters, we construct a *discriminant function* which gives us a decision boundary separating the two classes (hopefully).

With p -dimensional data, for each of the p dimensions we must estimate the *mean* of both distributions. In addition, we have *covariance* terms to estimate: for each distribution, we must estimate the covariance of the first dimension with the second, the covariance of the first with the third, the covariance of the second with the third, and so on and so forth.

- In p -dimensional space, how many parameters do we have to estimate for a generative model positing that our two classes of data are generated by multivariate normal distributions?

Overview of derivation

The most important information to take away from the derivations of discriminant analysis is the *set of assumptions* which underlies each method.

For binary classification, we first assume that the data from each class is drawn from a multivariate normal distribution. From this assumption alone, we can predict class membership from looking at whether or not the logarithm of the ratio of the likelihoods exceeds a certain threshold.

After simplification, we reduce to the following criterion for membership in class +1 (as opposed to class -1) where μ , Σ , and T represent distribution means, covariance matrices, and the threshold value:

$$\left[(\mathbf{x} - \mu_{+1})^T \Sigma_{+1}^{-1} (\mathbf{x} - \mu_{+1}) - (\mathbf{x} - \mu_{-1})^T \Sigma_{-1}^{-1} (\mathbf{x} - \mu_{-1}) \right] + \left[\ln \frac{|\Sigma_{+1}|}{|\Sigma_{-1}|} \right] > T.$$

There are quadratic terms on the left-hand side of the inequality, with entries of \mathbf{x} being multiplied with other entries of \mathbf{x} , so this is called **quadratic discriminant analysis**.

We can further simplify this expression by making the assumption that the two distributions have the same covariance matrix. In this case, the above criterion reduces to the form

$$\mathbf{w} \cdot \mathbf{x} > c$$

for a vector of weights \mathbf{w} and a constant threshold c which can be calculated from the parameters of the multivariate normal distributions.

Note that the left-hand side of the inequality is linear with respect to \mathbf{x} . As such, we call this **linear discriminant analysis**.

Feel free to skip the following section on the mathematical derivation of QDA and LDA if you find it too technically demanding.

Mathematical derivation

Suppose that we have p -dimensional data points \mathbf{x}_i , each with a binary class label $y_i \in -1, +1$. We make the assumption that the probability distribution for each class, $P(\mathbf{x} | y = -1)$ and $P(\mathbf{x} | y = +1)$, are both normally distributed with means and covariances (μ_{-1}, Σ_{-1}) and (μ_{+1}, Σ_{+1}) respectively.

With this assumption alone, we are ready to look at the *log-likelihood* of class membership. Specifically, we can predict class membership to be +1 if the log of the ratio of the likelihoods is above some threshold T , that is,

$$\log \frac{P(\mathbf{x} | y = +1)}{P(\mathbf{x} | y = -1)} > T.$$

Since we assumed that both probability densities are multivariate normal distributions, they have the functional form

$$P(\mathbf{x} | y = +1) = \frac{1}{(2\pi)^{p/2} |\Sigma_{+1}|^{1/2}} \exp \left(-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_{+1})^T \Sigma_{+1}^{-1} (\mathbf{x} - \boldsymbol{\mu}_{+1}) \right).$$

As such, after taking the logs and simplifying, we have the following criterion for membership in class +1:

$$\left[(\mathbf{x} - \boldsymbol{\mu}_{+1})^T \Sigma_{+1}^{-1} (\mathbf{x} - \boldsymbol{\mu}_{+1}) - (\mathbf{x} - \boldsymbol{\mu}_{-1})^T \Sigma_{-1}^{-1} (\mathbf{x} - \boldsymbol{\mu}_{-1}) \right] + \left[\ln \frac{|\Sigma_{+1}|}{|\Sigma_{-1}|} \right] > T.$$

As it stands, the classification of points with the above criterion is called **quadratic discriminant analysis** (QDA), because of the **cross terms** in the multiplication which lead to **quadric decision boundaries**.

We can *further* assume that the covariances of the two distributions are equal,¹ so $\Sigma = \Sigma_{-1} = \Sigma_{+1}$. After further simplification, we obtain the criterion

$$\mathbf{w} \cdot \mathbf{x} > c,$$

where

$$\mathbf{w} = \Sigma^{-1} (\boldsymbol{\mu}_{-1} - \boldsymbol{\mu}_{+1})$$

and

$$c = \frac{1}{2} \left(T - \boldsymbol{\mu}_{+1}^T \Sigma^{-1} \boldsymbol{\mu}_{+1} + \boldsymbol{\mu}_{-1}^T \Sigma^{-1} \boldsymbol{\mu}_{-1} \right).$$

This time, our solution is *linear* in \mathbf{x} , and accordingly is named **linear discriminant analysis** (LDA).

Note that the criterion $\mathbf{w} \cdot \mathbf{x} > c$ is equivalent to specifying that \mathbf{x} must be on a particular side of a *hyperplane* parameterized by \mathbf{w} and c . We can see this by realizing that the dot product $\mathbf{w} \cdot \mathbf{x}$ is proportional to the **projection** of vector \mathbf{x} onto vector \mathbf{w} . As such, if the dot product is positive, then \mathbf{x} is on one side of a

¹We also assume that the covariance matrices have full rank.

hyperplane orthogonal to \mathbf{w} , and if it negative, then \mathbf{x} is on the other side of the hyperplane.

Effectively, LDA finds a *separating hyperplane* which divides the two classes of training data, with slope and position set by \mathbf{w} and c .

Estimating parameters

The QDA and LDA criterion depend on knowing the covariances Σ_{-1} and Σ_{+1} (or just Σ) as well as the means μ_{-1} and μ_{+1} of the two multivariate distributions which we assume generate our data.

In practice, we do not actually know the true values of these parameters, but we can estimate them using the [sample mean and covariance](#) of our training data.

There is no general rule for choosing the value of the threshold T . If the data are projected onto \mathbf{w} , the one-dimensional distribution of projections can be analyzed to choose the threshold which seems to best separate the data.

Intuition

We will consider just LDA for simplicity, but the intuition here can be extended to QDA as well.

We know that LDA finds a hyperplane which is supposed to separate the two classes of data. However, what can we say about the properties of this hyperplane?

Since LDA works with the assumption that the two distributions (one corresponding to each class) are multivariate normal distributions, it effectively finds a separating hyperplane which does the following:

1. Maximizes the between-class variation, and
2. Minimizes the within-class variation.

“Variation” in the above refers to the notion of [Mahalanobis distance](#), which measures the distance between a point and distribution.

Because of the generative nature of the model, LDA works well when the underlying data really are drawn from multivariate normal distributions. Moreover, since the *between-class variation* is implicitly minimized by LDA, the two classes are well separated. However, if the means of the distributions are close together, LDA will struggle to find a good separating hyperplane.

Moreover, QDA needs to estimate many more parameters than LDA, so it often requires more data. If the covariances of the underlying distributions are very dissimilar, then QDA will eventually perform better than LDA; however, in

cases where the covariance matrices are relatively similar, it is possible that LDA will perform better than QDA up until a very high sample size is achieved.

Discriminant analysis in R

LDA and QDA are implemented in R as `lda()` and `qda()` in the MASS package. The resulting fit objects can be directly printed in the console for an overview of the calculated parameters, and class predictions can be made with `predict()` as usual. (For a prediction object, the actual class level assignments will be stored in `$class`.)

We'll first look at simulated data, where we know what the underlying distributions of the two classes look like.

- Generate 100 data points in the unit square. 50 of them should be drawn from a multivariate normal distribution centered at $(0.25, 0.25)$, and 50 of them should be drawn from a multivariate normal distribution centered at $(0.75, 0.75)$. Their covariance matrices should be identical, with the variance in each direction equal to 0.2 and the covariance between the dimensions equal to 0.1. Plot the data.
- Generate a vector with class labels for the data (-1 or +1 depending on which distribution a point is drawn from). Run both LDA and QDA on the data. Look at the estimates of the sample means and covariances and evaluate the accuracy of the predictions.
- Generate equal amounts of data from two multivariate normal distributions in the unit square with extremely different covariance matrices. Run both LDA and QDA on the data. How does the performance of each algorithm in classifying the training data vary as you change the size of the dataset?
- Generate 200 data points, 100 of which fall above the line $y = 1.5x + 0.2$ and 100 of which fall below the line $y = 1.5x + 0.05$. Plot the data.
- Run both LDA and QDA on the generated data and evaluate their performance. Do the same for smaller and larger datasets generated identically.

Next, we'll run our discriminant analysis methods on the aggregated speed dating dataset.

- Load the aggregated speed dating dataset and restrict to self-rated activity participation and the gender of the person rated.
- For each gender, compute the means and covariance matrices of the associated subset of the dataset. Compare them and predict the relative performance of LDA and QDA on this dataset.
- Create random subsets of the data that are 1%, 2%, 5%, 10%, 25%, 50%, 75%, and 100% of the size of the original dataset.

- Run LDA and QDA on each of the variably sized subsets of data, classifying gender with respect to self-rated activity participation. Plot their classification accuracy on the training data as a function of the proportion of the dataset used. Interpret the results.

Visualizing the results

Visualizing the results of LDA is a little troublesome. We will not attempt it here, but:

- Read [this StackExchange answer by “amoeba”](#) about visualizing the decision boundaries for *multiclass* LDA (briefly discussed below).

Extensions of discriminant analysis

LDA and QDA can be extended for usage in multiclass analysis. In addition, a variant of LDA is occasionally used for dimensionality reduction.

Logistic regression

You already have substantial experience applying logistic regression in a computational sense; as such, there will be no R exercises in this section. This section has two goals: first, to demonstrate further how logistic regression fits into the framework of classification and machine learning as a whole; second, to provide more refinement to your intuitions about how logistic regression works.

Like discriminant analysis, logistic regression is a *generative* method, which makes certain assumptions about the distributions from which the data are drawn. Instead of beginning with linear regression as motivation and developing logistic regression, we will instead begin from the standpoint of trying to answer the question: “How can we improve linear discriminant analysis?” Beautifully, we will see that the answers coincide.

For a review of the standard exposition of logistic regression, see Chapter 11 in [Advanced Data Analysis from an Elementary Point of View](#).

Mathematical derivation

The difficulty of the following mathematical exposition will vary from paragraph to paragraph. If you get caught up on understanding a specific statement, it’s better to move on and return to it later.

I had the pleasure of learning this derivation from *Statistical Methods for Pharmaceutical Research Planning* by Bergman and Gittins (1985).

The motivation behind the following derivation of logistic regression is the consideration of instances where the two classes in binary classification are *not* sampled from multivariate normal distributions with equal covariances. That is, we would like to relax the assumptions of LDA to allow for probability densities that are not normal.

In particular, instead of assuming that $P(\mathbf{x} \mid y = i) \propto \mathcal{N}(\boldsymbol{\mu}_i, \Sigma)$ for $i \in -1, +1$, where $\mathcal{N}(\boldsymbol{\mu}_i, \Sigma)$ denotes a multivariate normal distribution with mean $\boldsymbol{\mu}_i$ and covariance matrix Σ_i , we posit that

$$P(\mathbf{x} \mid y = i) \propto \mathcal{N}(\boldsymbol{\mu}_i, \Sigma) \phi(\mathbf{x}).$$

We relax our assumptions by saying that each of the two classes is drawn from a distribution equal to the *product* of a multivariate normal distribution with some unknown function $\phi(\mathbf{x})$.

The *discriminant functions* – that is, the essential inequalities which define QDA and LDA – stay the same! This is because in their derivations we take the *ratio* of the probability densities for the two classes, so the function $\phi(\mathbf{x})$ cancels and disappears. However, this is *not* a free lunch. A more subtle problem arises: the sample means and covariances, which we originally used as approximations to the parameters of the underlying multivariate normal distributions, are no longer the maximum likelihood estimators for those parameters in this more complex situation where we multiply the multivariate normal distribution by $\phi(\mathbf{x})$ (unless $\phi(\mathbf{x}) = 1$).

We will pause for a brief interlude in order to consider the *likelihood function* \mathcal{L} directly. The function $\mathcal{L}(\mathbf{x}, \boldsymbol{\theta})$ with respect to a collection of parameters $\boldsymbol{\theta}$ is equal to the *probability of observing the data \mathbf{x} given the parameter values $\boldsymbol{\theta}$ for the underlying distributions from which the data are generated*. To do *maximum likelihood estimation* of the parameters, we find the combination of parameters which maximizes \mathcal{L} .

Now, let us take write the training data as $T = T_{-1} \cup T_{+1}$, that is, the union of the classes -1 and $+1$. The likelihood function for the generative model we are considering is given by

$$\mathcal{L} = \prod_{\mathbf{x} \in T_{-1}} [P(\mathbf{x} \mid y = -1)P(y = -1)] \prod_{\mathbf{x} \in T_{+1}} [P(\mathbf{x} \mid y = +1)P(y = +1)].$$

Again, we can intuitively think of this as being the *probability of observing the data which we observed given the parameters of the underlying distributions*, where the parameters are implicitly encoded in the probabilities $P(y = -1)$ and $P(y = +1)$.

Expanding out the probabilities, we obtain

$$\mathcal{L} = \prod_{\mathbf{x} \in T_{-1}} [\gamma_{-1} P(y = -1) \mathcal{N}(\boldsymbol{\mu}_{-1}, \Sigma) \phi(\mathbf{x})] \prod_{\mathbf{x} \in T_{+1}} [P(y = +1) \mathcal{N}(\boldsymbol{\mu}_{+1}, \Sigma) \phi(\mathbf{x})]$$

by substituting in the explicit expressions for the distributions of each class, where the constants γ_{-1} and γ_{+1} are chosen so that the distributions properly integrate to 1. To extend the method of discriminant analysis even further, in the case where the function $\phi(\mathbf{x})$ is known we can maximize the likelihood function above and obtain the parameters of the multivariate normal distributions. However, the difficulty of this will vary depending on the form of $\phi(\mathbf{x})$.

We have more or less run into the limits of how far discriminant analysis can take us without using substantially more sophisticated methods. However, we can develop the method of *logistic classification* by first rewriting the likelihood function as follows:

$$\mathcal{L} = \prod_{\mathbf{x} \in T} P(\mathbf{x}) \left[\prod_{\mathbf{x} \in T_{-1}} P(y = -1 | \mathbf{x}) \prod_{\mathbf{x} \in T_{+1}} P(y = +1 | \mathbf{x}) \right].$$

What we have done here is *reverse the direction of the conditional dependences*. Instead of conditioning on class membership, we now look at the probability of class membership conditional on the location of each data point.

Now, due to laws of probability, we have the relation

$$P(\mathbf{x}) = P(\mathbf{x} | y = -1)P(y = -1) + P(\mathbf{x} | y = +1)P(y = +1).$$

Moreover, from Bayes' theorem, we have

$$P(y = +1 | \mathbf{x}) = \frac{P(\mathbf{x} | y = +1)P(y = +1)}{P(\mathbf{x} | y = -1)P(y = -1) + P(\mathbf{x} | y = +1)P(y = +1)}.$$

Substituting in the explicit expression for the probability densities from which each class is drawn and simplifying, we arrive at

$$P(y = +1 | \mathbf{x}) = \frac{\exp(\mathbf{x}^\top \mathbf{c} + d)}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)}$$

where $\mathbf{c} = \Sigma^{-1}(\boldsymbol{\mu}_{+1} - \boldsymbol{\mu}_{-1})$ and

$$d = -\frac{1}{2}(\boldsymbol{\mu}_{+1} - \boldsymbol{\mu}_{-1})^\top \Sigma^{-1}(\boldsymbol{\mu}_{+1} + \boldsymbol{\mu}_{-1}) + \log \frac{P(y = +1)\gamma_{+1}}{P(y = -1)\gamma_{-1}}.$$

Substituting the expressions for $P(\mathbf{x})$ and $P(y = +1 \mid \mathbf{x})$ into the likelihood function, we obtain

$$\mathcal{L} = l_{+1} l_{-1} \prod_{\mathbf{x} \in T} \{ \phi(\mathbf{x}) [\gamma_{+1} P(y = +1) \exp f(+1) + \gamma_{-1} P(y = -1) \exp f(-1)] \}$$

where

$$f(i) = -\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu}_i)^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}_i)$$

and

$$l_{+1} = \prod_{\mathbf{x} \in T_{+1}} \frac{\exp(\mathbf{x}^\top \mathbf{c} + d)}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)}$$

and

$$l_{-1} = \prod_{\mathbf{x} \in T_{-1}} \frac{1}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)}.$$

We can maximize this expression with respect to $\phi(\mathbf{x})$ as well as the means and shared covariance of the multivariate normal distributions. It is argued in Day and Kerridge (1967), *A General Maximum Likelihood Discriminant*, that the likelihood \mathcal{L} is primarily dependent on the values of \mathbf{c} and d , being relatively insensitive to the product over T in practice. As such, we obtain the following simplification for the *reduced likelihood*:

$$\mathcal{L}_{\text{red}} = l_{+1} l_{-1} = \prod_{\mathbf{x} \in T_{+1}} \frac{\exp(\mathbf{x}^\top \mathbf{c} + d)}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)} \prod_{\mathbf{x} \in T_{-1}} \frac{1}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)}.$$

Our task is much simpler now: to simply maximize this straightforward-seeming function with respect to the vector \mathbf{c} and the constant d . Equivalently, we can recast the problem into an even simpler form by taking the *logarithm*, which turns products into sums:

$$\log \mathcal{L}_{\text{red}} = \sum_{\mathbf{x} \in T} -\log(1 + \exp(\mathbf{x}^\top \mathbf{x} + d)) + \sum_{\mathbf{x} \in T_{+1}} (\mathbf{x}^\top \mathbf{c} + d).$$

Looking at Shalizi's *Advanced Data Analysis from an Elementary Point of View*, equation 11.10 for the log-likelihood in his derivation of logistic regression is:

$$\mathcal{L}(\beta_0, \beta) = \sum_{i=1}^n -\log(1 + e^{\beta_0 + \mathbf{x}_i \cdot \beta}) + \sum_{i=1}^n y_i (\beta_0 + \mathbf{x}_i \cdot \beta).$$

This is the same as our own log-likelihood $\log \mathcal{L}_{\text{red}}$! The only difference in the second sum can be explained by the fact that in Shalizi's derivation, we have $y_i \in \{0, 1\}$ for the two classes, so in effect the sum is only over the values of $\beta_0 + \mathbf{x}_i \cdot \beta$ for $x_i \in \{x_i \mid y_i = 1\}$.

Indeed, Bergman and Gittins themselves make the following remark:

The logistic classification procedure may also be derived by a different line of argument. One could assume that the underlying population has a logistic density with unknown parameters \mathbf{c} and d . This again yields the classification function $\mathbf{x}^\top \mathbf{c} + d$. This is of some interest, for it implies that classification on the basis of the model above is essentially the same as classification on the basis of the log-linear model with no interaction terms (see, *e.g.*, Bishop *et al.*, 1975).

For data in p -dimensional space, we only have $p + 1$ parameters to estimate! This can easily be done using a standard optimization technique such as the [Newton-Raphson method](#). After estimating the parameters, we can then explicitly calculate the probability of class membership for any data point with

$$P(y = +1 \mid \mathbf{x}) = \frac{\exp(\mathbf{x}^\top \mathbf{c} + d)}{1 + \exp(\mathbf{x}^\top \mathbf{c} + d)}$$

and, of course, $P(y = -1 \mid \mathbf{x}) = 1 - P(y = +1 \mid \mathbf{x})$. Finally, we can classify a point \mathbf{x} as belonging to group $+1$ if and only if $P(y = +1 \mid \mathbf{x}) > T$ for some threshold T depending on the costs of misclassification for each class.

Guaranteed linear separation

Interestingly, the method of logistic regression is guaranteed to return a hyperplane which perfectly separates the two classes in the training data if indeed it is possible to do so.

Indeed, suppose that $g(\mathbf{x}) = \mathbf{x}^\top \mathbf{c} + d$ be an arbitrary linear classification function, and assume without loss of generality that $g(\mathbf{x})$ classifies \mathbf{x} as belonging to class $+1$ if and only if $g(\mathbf{x}) > 0$. Now, the corresponding reduced likelihood function is given by

$$\mathcal{L}_{\text{red}}(g(\mathbf{x})) = \prod_{\mathbf{x} \in T_{+1}} \frac{\exp g(\mathbf{x})}{1 + \exp g(\mathbf{x})} \prod_{\mathbf{x} \in T_{-1}} \frac{\exp -g(\mathbf{x})}{1 + \exp -g(\mathbf{x})},$$

since

$$\frac{\exp -g(\mathbf{x})}{1 + \exp -g(\mathbf{x})} = \frac{1}{1 + \exp g(\mathbf{x})}.$$

From the form of the product, we may see that each term takes on a value in $[1/2, 1]$ if the corresponding point \mathbf{x} is correctly classified and a value in $[0, 1/2]$ otherwise.

Now, assume that $g^*(\mathbf{x})$ is a linear classification function which correctly classifies all $\mathbf{x} \in T$, that is to say, $g^*(\mathbf{x}) > 0$ if $\mathbf{x} \in T_{+1}$ and $g^*(\mathbf{x}) < 0$ if $\mathbf{x} \in T_{-1}$. If $g^*(\mathbf{x})$ is a perfect classification function, then so is $kg^*(\mathbf{x})$ for any constant $k > 0$, so

$$\lim_{k \rightarrow \infty} \mathcal{L}_{\text{red}}(kg^*(\mathbf{x})) = 1.$$

Since the function $\mathcal{L}_{\text{red}}(kg^*(\mathbf{x}))$ is continuous with respect to k , it must be the case that there exists some value $k = k^*$ such that

$$\mathcal{L}_{\text{red}}(k^*g^*(\mathbf{x})) > \frac{1}{2}.$$

In other words, there exists a function $g'(\mathbf{x}) = k^*g^*(\mathbf{x})$ such that $\mathcal{L}_{\text{red}}(g'(\mathbf{x})) > 1/2$. Therefore, as we maximize \mathcal{L}_{red} , we will eventually arrive at values of the parameters \mathbf{c} and \mathbf{d} which correspond to the function $g'(\mathbf{x})$, which by construction fully separates the two classes. (Indeed, if $g'(\mathbf{x})$ were *not* fully separating, then one of the terms in the product $\mathcal{L}_{\text{red}}(g'(\mathbf{x}))$ would be below $1/2$, and so the entire product would also be below $1/2$.)

Closing notes

Given the popularity of logistic regression today, it is amusing that Bergman and Gittins say:

Apparently, this model has yet to receive attention in the pharmaceutical literature. In view of its promise, it certainly warrants some consideration

Perceptrons

We will proceed to considering the **perceptron algorithm**, one of the first and simplest purely predictive classification techniques developed. Unlike the methods of discriminant analysis and logistic regression, which are *generative* meth-

ods, the perceptron simply aims to find the best hyperplane possible for class separation.

- Read about the history of the perceptron on [Wikipedia](#).

For labeled training data in n -dimensional space, the perceptron algorithm will attempt to construct an $(n - 1)$ -dimensional hyperplane which separates the two classes.

Getting started

First, we will generate linearly separable data.

- Generate 1000 data points falling above the line $1.5x + 0.2$ and 1000 data points falling below the line $1.5x + 0.05$, binding them into the same matrix.
- Create a `labels` vector where the i th entry is 1 or -1 according to whether the i th data point in the matrix was generated with `label=1` or `label=-1`. Graph the results with `qplot()`, using the `color=labels` parameter to color the points.
- Add a column to your matrix of data consisting purely of 1s. The perceptron algorithm needs this to work properly; it serves as an “intercept term” that allows the separating hyperplane it generates to be shifted up or down instead of being necessarily centered at 0.

More precisely: for each point x_i , the perceptron algorithm forms an augmented point $x'_i = (x, 1)$. It works by updating a vector of *weights* w step-by-step as it iterates over the augmented data points, adjusting the weights whenever it encounters a training point which it misclassifies. To perform classification, it looks at the *sign* of $w \cdot x'_i$, classifying positive results as 1 and negative results as -1.

- Write a function `dot(x, y)` which computes the [dot product](#) between the vectors x and y , denoted as $x \cdot y$. It is equal to the sum of $x_1y_1 + x_2y_2 + \dots + x_ny_n$.
- Write a function `perceptron(xs, y, w, rate, niter)` following these specifications:
 - Take as input a matrix of data `xs`, a vector of labels (1 or -1) `y`, a vector of weights `w`, a learning rate `rate`, and a number of iterations `niter`. `niter` should default to the number of data points.
 - Generate `niter` random indices from the row indices of the data, sampled *without* replacement.
 - For each sampled row, do the following: For convenience, call the sampled row `xi`. Compute the *dot product* between `xi` and `w`. The

classification of x_i , according to the perceptron, is the `sign()` of the dot product. If the classification is a false negative, update w by adding $\text{rate} \cdot x_i$, and if the classification is a false positive, update w by subtracting $\text{rate} \cdot x_i$.

- Return the weights after the `niter` iterations have finished.

Visualizing your results

The weights of the perceptron parametrize a *line* given by $w[1] \cdot x + w[2] \cdot y + w[3] = 0$. This is called the *decision boundary*. Intuitively, the vector w , if we ignore its last entry, is a vector which is *perpendicular* to the separating hyperplane.

- What are the slope and y -intercept of this line in terms of the components of w ?
- Set the seed for consistency and try `perceptron()` on your generated data with `rate=1` and the default value for `niter`. You should initialize the weights vector to a vector of zeroes and then pass in the output of the first `perceptron()` call to the initial weights of the next. After each call to `perceptron()`, plot the *decision boundary* corresponding to your results overlaid on top of the scatterplot of data points by adding a `geom_abline()` call with the intercept and slope parameters.

A perceptron is *guaranteed to converge* to a line which fully separates two class *if* the two classes are linearly separable. If a perceptron converges to a solution, that means that it will correctly classify all of the training data.

- Write a function `perceptron_conv()` which takes in all the arguments of `perceptron()` aside from `niter`. It should run `perceptron()` until the solution converges. It should return a list with the final weights as well as the number of total iterations (it's fine if the number of iterations is an overestimate).
- Set the seed for consistency and plot the decision boundary which `perceptron_conv()` gives you for your output. Try a variety of different seeds. How much variation do you observe?
- How does the learning rate of the perceptron affect the speed of convergence for your current data? What if you have 20 data points (generated in the same way) instead of 2000?

Congratulations: you have successfully implemented one of the classic algorithms of machine learning. It even counts as a neural network (with a single neuron)!

Regarding the perceptron's advantages, Vapnik writes:

By the time the Perceptron was introduced, classical discriminant analysis based on Gaussian distribution functions had been studied in great detail. [...]

[T]o construct this model using classical methods requires the estimation of about $0.5n^2$ parameters where n is the dimensionality of the space. Roughly speaking, to estimate one parameter of the model requires C examples. Therefore to solve the ten-digit recognition problem using the classical technique one needs $\approx 10(400)^2C$ examples. The Perceptron used only 512.

This shocked theorists. It looked as if the classical statistical approach failed to overcome the curse of dimensionality in a situation where a heuristic method that minimized the empirical loss easily overcame this curse.

Unfortunately, despite its advantages, the standard perceptron cannot classify data which is not linearly separable, because the algorithm will simply not converge.

Also, as you have seen, when it *does* converge, the perceptron's solution is highly dependent upon choice of random seed. When the training data are linearly separable *and* there exists a "gap" or "margin" between them, there are infinitely many possible choices of a valid decision boundary. We will see later how to resolve these problems.

- Read more about the fascinating history of perceptrons in the Wikipedia article for Minsky and Papert's [book about perceptrons](#).

Support vector machines

Kernel methods

Decision trees for classification

Closing notes